

Binary

- > Thresholding
- > Threshold detection
- > Variations
- > Mathematical Morphology

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 1

1

Thresholding

- Distinct foreground & background needed

Original Threshold = 10

How do we determine the best threshold?

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 3

3

Thresholding

- Binary thresholding
 - for all pixels
 $g(i,j) = 1 \text{ for } f(i,j) \geq T$
 $= 0 \text{ for } f(i,j) < T$
 - Simple scenes?
 - LUT
 - for all grey levels
 $LUT(k) = 1 \text{ for } k \geq T$
 $= 0 \text{ for } k < T$
 - for all pixels
 $g(i,j) = LUT(f(i,j))$
 - Objects of interest vs. background

```
threshold(gray_image,binary_image,threshold,
          255,THRESH_BINARY);
```

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 2

2

Threshold Detection

- Manual Setting
 - Issue of changing lighting
- Need to determine automatically

For the techniques which follow:

- Image – $f(i,j)$
- Histogram – $h(g)$
- Probability Distribution – $p(g) = h(g) / \sum_g h(g)$

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 4

4

1

2

Threshold Detection – Otsu Thresholding

- Minimize the spread of the pixels...

- Smallest within class variance

$$\sigma_w^2(T) = w_f(T)\sigma_f^2(T) + w_b(T)\sigma_b^2(T)$$

$$w_f(T) = \sum_{g=0}^{255} p(g) \quad \sigma_f^2(T) = \frac{\sum_{g=0}^{255} p(g) \cdot (g - \mu_f(T))^2}{w_f(T)}$$

$$w_b(T) = \sum_{g=0}^{T-1} p(g) \quad \sigma_b^2(T) = \frac{\sum_{g=0}^{T-1} p(g) \cdot (g - \mu_b(T))^2}{w_b(T)}$$

$$\mu_f(T) = \frac{\sum_{g=0}^{255} p(g) \cdot g}{w_f(T)} \quad \mu_b(T) = \frac{\sum_{g=0}^{T-1} p(g) \cdot g}{w_b(T)}$$

- Largest between class variance

$$\sigma_B^2(T) = w_f(T)w_b(T)(\mu_f(T) - \mu_b(T))^2$$

Binary

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 5

Variations – Adaptive Thresholding

- The adaptive thresholding algorithm is

- Divide the image into sub-images,
- Compute thresholds for all sub-images,
- Interpolate thresholds for every point using bilinear interpolation.

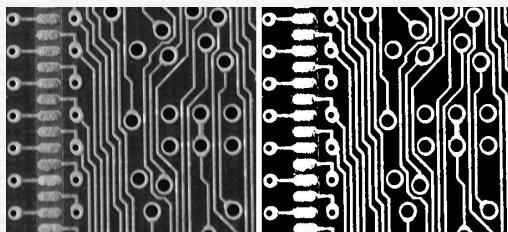


Binary

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 7

Threshold Detection – Otsu Thresholding



```
threshold( gray_image, binary_image, threshold,
           255, THRESH_BINARY | THRESH_OTSU );
```

Binary

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 6

Variations – Adaptive Thresholding

- OpenCV version:

```
if ((f(i,j) - (\sum_{a=-m...m, b=-m...m} f(i+a,j+b) / (2m+1)^2)) > offset)
    g(i,j) = 255
else g(i,j) = 0
```

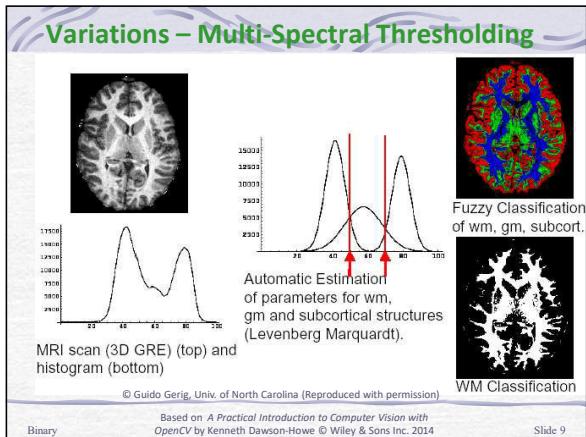


```
adaptiveThreshold( gray_image,binary_image,output_value,
                   ADAPTIVE_THRESH_MEAN_C,THRESH_BINARY,
                   block_size,offset );
```

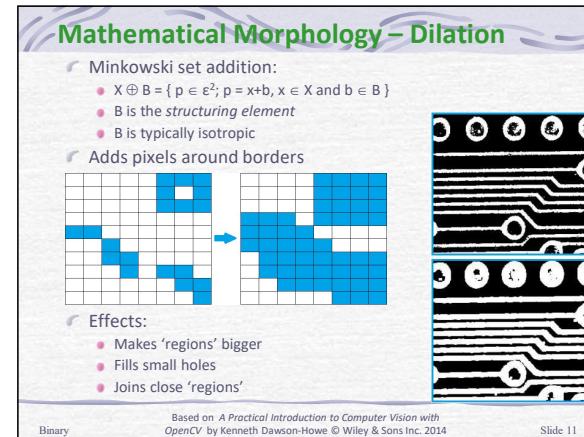
Binary

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

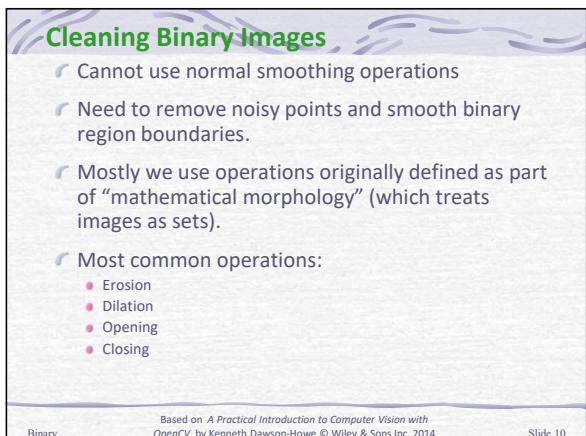
Slide 8



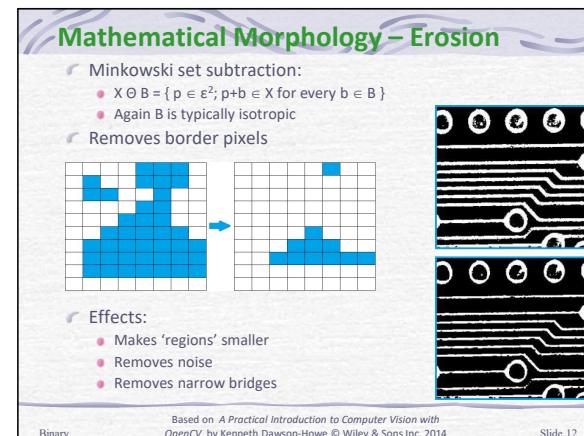
9



11



10



12

Mathematical Morphology –

- Opening:** $X \circ D = (X \ominus D) \oplus D$
 - Removes noise
 - Removes narrow bridges
 - Roughly maintains 'region' size
 - Smooths shape
- Closing:** $X \bullet D = (X \oplus D) \ominus D$
 - Fills small holes
 - Joins close 'regions'
 - Roughly maintains 'region' size

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Binary Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014 Slide 13

13

Mathematical Morphology – OpenCV Code

```
dilate( binary_image, dilated_image, Mat() );
Mat structuring_element( 5, 5, CV_8U, Scalar(1) );
dilate( binary_image, dilated_image, structuring_element);

erode( binary_image, eroded_image, Mat() );
Mat structuring_element( 5, 5, CV_8U, Scalar(1) );
erode( binary_image, eroded_image, structuring_element);

Mat five_by_five_element( 5, 5, CV_8U, Scalar(1) );
morphologyEx( binary_image, opened_image,
    MORPH_OPEN, five_by_five_element );
morphologyEx( binary_image, closed_image,
    MORPH_CLOSE, five_by_five_element );
```

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Binary Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014 Slide 15

15

Mathematical Morphology –

- Opening:** $X \circ D = (X \ominus D) \oplus D$
- Closing:** $X \bullet D = (X \oplus D) \ominus D$
- Properties**
 - $X \circ D = (X \ominus D) \oplus D$ and $X \bullet D = (X \bullet D) \ominus D$

Isotropic structuring element:

- Eliminates small image details

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Binary Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014 Slide 14

14

Mathematical Morphology – Greyscale / Colour

One set per grey level (g)

- All points $\geq g$...

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Binary Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014 Slide 16

16

Mathematical Morphology – Local maxima

- Can be used to locate local maxima and minima

```
Mat dilated, thresholded_input, local_maxima, thresholded_8bit;  
dilate( input, dilated, Mat() );  
compare( input, dilated, local_maxima, CMP_EQ );  
threshold( input, thresholded_input, threshold, 255,  
          THRESH_BINARY );  
thresholded_input.convertTo( thresholded_8bit, CV_8U );  
bitwise_and( local_maxima, thresholded_8bit, local_maxima );
```

Binary

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 17

Edges

- Edge Detection
- Contour Segmentation
- Hough Transform
- Least Squared Error
- Random Sample Consensus (RANSAC)

Edges

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 1

Edges

- An approach to segmentation.
- The analysis of the discontinuities in an image.



- No correct answer?
- An alternative to region based processing.

Edges

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 2

Edge Detection – Topics

- ✓ 1st derivative edge detection
- ✓ 2nd derivative edge detection
- ✓ Multispectral edge detection
- ✓ Image sharpening

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Edges Slide 3

3

Edge Detection – 1st derivative definitions



Calculus...

- Rate of change in two directions
- Vector variable:
 - Gradient Magnitude
 - Orientation (0° is East)

$$\nabla f(i,j) = \sqrt{\left(\frac{\delta f(i,j)}{\delta i}\right)^2 + \left(\frac{\delta f(i,j)}{\delta j}\right)^2}$$

$$\phi(i,j) = \arctan\left(\frac{\delta f(i,j)}{\delta j}, \frac{\delta f(i,j)}{\delta i}\right)$$

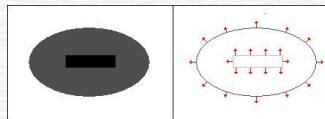
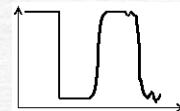
Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Edges Slide 5

5

Edge Detection – What is an edge?

- ✓ Where brightness changes abruptly
- ✓ Edges have
 - Magnitude (Gradient)
 - Direction (Orientation)
- ✓ Edge Profiles
 - Step
 - Real
 - Noisy

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Edges Slide 4

4

Edge detection – Digital images

- ✓ Derivatives work on continuous functions
 - Map every point in the input image to the output
- ✓ Discrete domain
 - Differences
 - Orthogonal

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Edges Slide 6

6

2

3

Edge detection – 1st derivative – Roberts

$\delta_1(i, j) = f(i, j) - f(i + 1, j + 1)$

$\delta_2(i, j) = f(i, j + 1) - f(i + 1, j)$

- Convolution Masks

$$h_1(i, j) = \begin{bmatrix} 1 & 0 \\ 0 & -1 \end{bmatrix} \quad h_2(i, j) = \begin{bmatrix} 0 & 1 \\ -1 & 0 \end{bmatrix}$$

- Sensitivity to Noise
- Binary Images

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Edges

Slide 7

7

Edge detection – 1st derivative – Compass

- Compass edge detectors:
- Partial derivatives defined for a number of orientations (typically 8)
- Prewitt:

$$h_1(i, j) = \begin{bmatrix} 1 & 0 & 1 \\ 0 & 1 & 0 \\ -1 & -1 & -1 \end{bmatrix} \quad h_2(i, j) = \begin{bmatrix} 0 & 1 & 1 \\ -1 & 0 & 1 \\ -1 & -1 & 0 \end{bmatrix} \quad h_3(i, j) = \begin{bmatrix} -1 & 0 & 1 \\ -1 & 0 & 1 \\ -1 & 0 & 1 \end{bmatrix} \quad h_4(i, j) = \begin{bmatrix} -1 & -1 & 0 \\ -1 & 0 & 1 \\ 0 & 1 & 1 \end{bmatrix}$$

$$h_5(i, j) = \begin{bmatrix} -1 & -1 & -1 \\ 0 & 0 & 0 \\ 1 & 1 & 1 \end{bmatrix} \quad h_6(i, j) = \begin{bmatrix} 0 & -1 & -1 \\ 1 & 0 & -1 \\ 1 & 1 & 0 \end{bmatrix} \quad h_7(i, j) = \begin{bmatrix} 1 & 0 & -1 \\ 1 & 0 & -1 \\ 1 & 0 & -1 \end{bmatrix} \quad h_8(i, j) = \begin{bmatrix} 1 & 1 & 0 \\ 1 & 0 & -1 \\ 0 & -1 & -1 \end{bmatrix}$$

- Use $h_3(i, j)$ and $h_1(i, j)$ to derive gradient and orientation

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Edges

Slide 9

9

Edge detection – Digital images

- Need to define the partial derivatives so that they:
 - Cross at a single middle point
 - Preferably cross at the centre of a pixel
 - Evaluate points which are not too close together
 - Deal with some degree of image noise

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Edges

Slide 8

8

Edge detection – 1st derivative – Compass

- Sobel

$$h_1(i, j) = \begin{bmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix} \quad h_3(i, j) = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}$$

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Edges

Slide 10

10

Edge detection – 1st derivative – Compass

$\nabla f(i,j) = \sqrt{\left(\frac{\delta f(i,j)}{\delta i}\right)^2 + \left(\frac{\delta f(i,j)}{\delta j}\right)^2}$ $\phi(i,j) = \arctan\left(\frac{\delta f(i,j)}{\delta j}, \frac{\delta f(i,j)}{\delta i}\right)$

$$\nabla f(i,j) = \left|\frac{\delta f(i,j)}{\delta i}\right| + \left|\frac{\delta f(i,j)}{\delta j}\right|$$

```
Mat horizontal_derivative, vertical_derivative;
Sobel( gray_image, horizontal_derivative, CV_32F,1,0 );
Sobel( gray_image, vertical_derivative, CV_32F,0,1 );
Mat abs_gradient, l2norm_gradient, orientation;
abs_gradient = abs(horizontal_derivative) +
    abs(vertical_derivative);
cartToPolar(horizontal_derivative,vertical_derivative,
    l2norm_gradient,orientation);
```

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Edges Slide 11

11

Edge detection – 1st derivative – Non maxima suppression

- Use orientation information
- Algorithm:
 - Quantise edge orientations
 - For all points (i,j)
 - Look at the 2 points orthogonal to edge
 - if $\text{gradient}(i,j) < \text{gradient}(\text{either of these 2 points})$
 - $\text{output}(i,j) = 0$
 - $\text{else } \text{output}(i,j) = \text{gradient}(i,j)$
 - Now thresholding can be used... or hysteresis thresholding (detailed later...)

Gradients Threshold=10 Edges

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Edges Slide 13

13

Edge detection – 1st derivative – Thresholding

- Simple thresholding
 - Too many points
 - Too few points

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Edges Slide 12

12

Edge detection – 1st derivative – NMS

```
nms_result = gradients.clone();
for (int row=1; row < gradients.rows-1; row++)
  for (int column=1; column < gradients.cols-1; column++)
  {
    float curr_gradient = gradients.at<float>(row,column);
    float curr_orientation = orientations.at<float>(row,column);
    // Determine which neighbours to check
    int direction = (((int)(16.0*(curr_orientation)/(2.0*PI))+15)%8)/2;
    float gradient1 = 0.0, gradient2 = 0.0;
    switch(direction)
    {
    case 0:
      gradient1 = gradients.at<float>(row-1,column-1);
      gradient2 = gradients.at<float>(row+1,column+1);
      break;
```

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Edges Slide 14

7

6

Edge detection – 1st derivative – NMS

```

case 1:
gradient1 = gradients.at<float>(row-1,column);
gradient2 = gradients.at<float>(row+1,column);
break;
case 2:
gradient1 = gradients.at<float>(row-1,column+1);
gradient2 = gradients.at<float>(row+1,column-1);
break;
case 3:
gradient1 = gradients.at<float>(row,column+1);
gradient2 = gradients.at<float>(row,column-1);
break;
}
if ((gradient1 > curr_gradient) || (gradient2 > curr_gradient))
nms_result.at<float>(row,column) = 0.0;
}

```

Edges

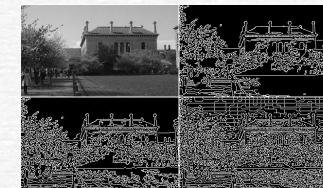
Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 15

15

Edge detection – Canny algorithm

1. Convolve image with Gaussian
2. Compute the first derivative gradients and orientations
3. Apply non-maxima suppression
4. Apply a double threshold to find strong & weak edge points.
5. Threshold edges with hysteresis
 - Remove any weak edges points not connected to strong edge points.



Edges

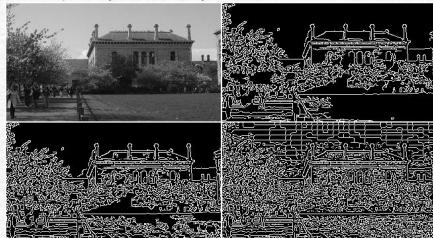
Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 17

17

Edge detection –Canny

- Combination of first and second derivative detectors
- Second derivative intended to improve localisation
- Thresholding incorporated so result is binary
- (Optional?) Analysis at multiple scales



Edges

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 16

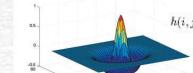
16

Edge detection – Canny – uses 2nd derivative?

- According to the original paper Canny is supposed to be a combination of first and second derivative detectors
- Second derivative intended to improve localisation



- Uses a filter like the Laplacian (of Gaussian)



$$h(i,j) = \begin{bmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{bmatrix}$$



- Need to find zero-crossings.

Edges

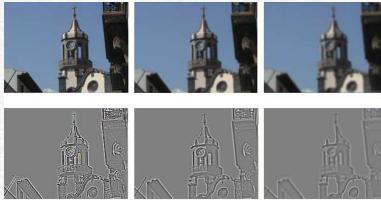
Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 18

18

Edge detection –Canny

- Also according to the original paper the image is supposed to be analysed at multiple scales
 - Smooth image using different Gaussians
 - Look for correspondences across scale to identify significant discontinuities



Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Edges Slide 19

19

Multispectral edge detection

- Grey levels are ordered, colours are not.
- Multispectral edge detection is addressed in 3 ways:
 - Output fusion
 - Combine edges
 - Multidimensional gradient methods
 - Combine gradients
 - Vector methods
 - Distance between colours in 3D



Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Edges Slide 21

21

Edge detection – advanced – Canny



```
Canny( gray_image, binary_edges, 100, 200 );
```

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Edges Slide 20

20

10

Contour Segmentation – Topics

- Edge data representations
- Border Detection
- Line segment extraction

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Edges Slide 22

11

Contour Segmentation – Boundary Chain Codes

- Each chain contains
 - Start point
 - A list of orientations

Features!

- Orientation & Scale dependent
- Slightly position dependent
- Can be smoothed to reduce boundary noise
- Difficulties obtaining consistent representation from image.

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Edges

Slide 23

23

Contour Segmentation – Boundary Chain Codes

```
vector<vector<Point>> contours;
vector<Vec4i> hierarchy;
findContours( binary_edge_image, contours, hierarchy,
              CV_RETR_CCOMP, CV_CHAIN_APPROX_NONE );

for (int contour_number=0;
     (contour_number<contours.size()); contour_number++)
{
    Scalar colour( rand()&0xFF, rand()&0xFF, rand()&0xFF );
    drawContours( display_image, contours, contour_number,
                  colour, 1, 8, hierarchy );
}
```

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Edges

Slide 25

25

Contour Segmentation – Boundary Chain Codes

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Edges

Slide 24

24

Contour Segmentation

Represent an edge image as a graph:

- Node n_i
 - Corresponds to pixel x_i
 - Associated cost $s(x_i)$ and orientation $\phi(x_i)$
- n_i, n_j are connected by an arc if
 - Pixels x_i, x_j are 8-connected neighbours
 - $\phi(x_i)$ and $\phi(x_j)$ match the local border:
 - x_j must be one of 3 neighbours in the direction $[\phi(x_i)-\pi/4, \phi(x_i)+\pi/4]$
 - $|\phi(x_i) - \phi(x_j)| < \pi/2$

When expanding edge chains consider

- Strength of edges
 - E.g. $(\max_{image} s(x_k)) - s(x_i)$
- Border curvature
 - $\text{diff}[\phi(x_i) - \phi(x_j)]$

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Edges

Slide 26

26

Contour Segmentation

Purpose: Segment all edge chains in an image

Algorithm:

1. Search for the strongest (unused) node in the graph. If none go to step 5.
2. Expand all the edges in front of the specified edge
3. Expand all the edges behind of the specified edge
4. If the edge chain consists of > 3 pixels store it. Go to step 1.
5. Modify edge chains to fill small breaks and to remove duplicate contours.
6. Repeat step 5 until stable.

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Edges Slide 27

27

Contour Segmentation – Straight line extraction example

Epsilon = 1

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Edges Slide 29

29

Contour Segmentation – Segment sequences

- Boundary
- Sequence of segments
- Segment
 - Start and end points: x_1 and x_2
 - Type: straight line, some type of curve, etc.
 - Other parameters
- Accurate polygonal representation
 - Must define acceptable tolerances
 - Vertices – where?

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Edges Slide 28

28

14

Contour Segmentation – Obtaining polygonal segments

- Recursive boundary splitting
 - Split at furthest point
 - Keep going until with tolerance

```
vector<vector<Point>> approx_contours(contours.size());
for (int contour_number=0;
     (contour_number<contours.size()); contour_number++)
    approxPolyDP( Mat(contours[contour_number]),
                  approx_contours[contour_number], 3, true );
```

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Edges Slide 30

30

15

Hough transform

- Direct transformation from image space to probability of the existence of some feature...
 - Lines
 - Circles
 - Generalised shapes

© John McDonald, National University of Ireland, Maynooth from J. McDonald, I. Franz and R. Shorten. Application of the Hough Transform to Lane Detection in Motorway Driving Scenarios. Published in Proc. of the Irish Signals and Systems Conference, 2001. R Shorten, T. Ward, T. Lysaght (Eds) (Reproduced with permission)

Based on A Practical Introduction to Computer Vision with OpenCV by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Edges Slide 31

31

Hough Transform – Line detection

Based on A Practical Introduction to Computer Vision with OpenCV by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Edges Slide 33

33

Hough Transform – Line detection

Line equation:

- $j = m \cdot i + c$
- Lines in Hough space
- What about $i = c$
- $r = i \cdot \cos \theta + j \cdot \sin \theta$
- Sinusoidal curves

Based on A Practical Introduction to Computer Vision with OpenCV by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Edges Slide 32

32

Hough Transform – Circle detection

Equation of a circle $(i-a)^2 + (j-b)^2 = r^2$

- Assume constant radius r

Transform

- From Image space (x, y)
- To Hough space (a, b)

Algorithm

- Initialise accumulator to 0
- For every edge point
 - Increment cells in accumulator corresponding to all possible circle centers
- Search for Maximums

Based on A Practical Introduction to Computer Vision with OpenCV by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Edges Slide 34

34

Hough Transform – Circle detection

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Edges

Slide 35

35

Hough Transform – Resolution of Hough Space

- Circle**
 - (a, b) could be anywhere up to r outside the image space
 - Can detect the circle centre to higher precision than the image points
- Lines**
 - $-\text{dist}(0,0, M,N) \leq r \leq +\text{dist}(0,0, M,N)$
 - $-\pi \leq \theta \leq \pi$
 - Precision of s and θ application dependent

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Edges

Slide 37

37

Hough transform – in OpenCV

- Hough for lines:

```
vector<Vec2f> hough_lines;
HoughLines( binary_edge_image, hough_lines, 1, PI/200.0, 60);
```
- Probabilistic Hough for line segments:

```
vector<Vec4i> line_segments;
HoughLinesP( binary_edge_image, line_segments, 1.0,
PI/200.0, 20, 20, 5);
```
- Hough for circles:

```
vector<Vec3f> circles;
HoughCircles( gray_image, circles, CV_HOUGH_GRADIENT,
2,20,300,20,5,15);
```

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Edges

Slide 36

36

Hough Transform – Efficient Implementation

- Use edge orientations
 - Restrict the mapping into space
- Use half the accumulator
 - For lines only
- Multi-resolution
 - Process at a small resolution
 - Higher resolution to find accurate data
- Problem
 - What if the size of the circle is unknown
 - 3-D Accumulator??

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Edges

Slide 38

38

18

19

Other techniques for estimating pose (lines ++)

- Common to have a number of data points which may
 - form part of a feature (such as a straight line) OR
 - represent possible matches with some known object features
- We need to determine the best possible position or pose of this feature or object
- Techniques:
 - Least squared error solution
 - RANSAC

Recognition

© Kenneth Dawson-Howe 2015

Slide 39

39

Least squared error solution

Given (x_i, y_i) where $i = 1..N$

$$\mu_x = \frac{1}{N} \sum_{i=1}^N x_i \quad \mu_y = \frac{1}{N} \sum_{i=1}^N y_i$$

$$\sigma_x = \sqrt{\frac{1}{N} \sum_{i=1}^N (x_i - \mu_x)^2} \quad \sigma_y = \sqrt{\frac{1}{N} \sum_{i=1}^N (y_i - \mu_y)^2}$$

$$\text{Pearson's correlation coefficient: } \rho_{xy} = \frac{\text{cov}_{xy}}{\sigma_x \sigma_y} = \frac{\sum_{i=1}^N (x_i - \mu_x)(y_i - \mu_y)}{\sigma_x \sigma_y}$$

$$m = \rho_{xy} \frac{\sigma_y}{\sigma_x}$$

$$c = \mu_y - m \cdot \mu_x$$

Recognition

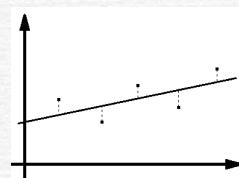
© Kenneth Dawson-Howe 2015

Slide 41

41

Least squared error solution

- Linear fit which best matches the data.
- For a straight line it minimises the sum of the vertical residuals



- Line equation: $y = m \cdot x + c$
- First compute the slope m
- Then the intercept c

Recognition

© Kenneth Dawson-Howe 2015

Slide 40

40

Least squared error solution – Problems

- Assumes line is not vertical
- Assumes that error distribution is normal
- Assumes that the points which should be included in the regression are known (i.e. segmented)
- Assumes no significant outliers

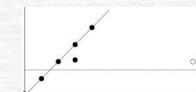


Figure 10.7: Influence of an outlier in least squares line fitting. With 6 valid data points and 1 gross outlier (white), the best line is shown in solid. Least squares, followed by discarding the worst outlier, reaches the dotted line after 3 discards [Fischler and Bolles, 1981]. © Cengage Learning 2015.
© Cengage Learning Engineering. Reproduced with permission. From *Image Processing, Analysis and Machine Vision* by Milan Sonka, Vaclav Hlavac and Roger Boyle

Recognition

© Kenneth Dawson-Howe 2015

Slide 42

42

20

21

Random Sample Consensus (RANSAC)

- Uses the minimum number of data points (m) to determine the model
 - For a straight line $m=2$

Technique

1. Randomly select the minimum number (m) of data points from the N data points (x_i, y_i) where $i = 1..N$
2. Determine the model from the selected data points
3. Determine how many data points are within some tolerance of the model – the consensus set
4. If the consensus set is not big enough (i.e. is smaller than some pre-set threshold) go back to step 1 (OR fail if tried too often).
5. If the consensus set was big enough, re-compute the model using all points in the consensus set.

Recognition © Kenneth Dawson-Howe 2015 Slide 43

43

RANSAC Example

Figure 10.8: Use of RANSAC in panoramic stitching (the original images are, of course, in color): (a) An image pair that overlaps. (b) Each dot represents an "interested point," whose vector matches a point in the other image. (c) The illustrated dots are the RANSAC "inliers." (d) Prior to smoothing, the resulting overlap. Courtesy of D. Lowe, M. Brown, University of British Columbia.

© Cengage Learning Engineering. Reproduced with permission. From *Image Processing, Analysis and Machine Vision* by Milan Sonka, Vaclav Hlavac and Roger Boyle

Recognition Slide 45

45

RANSAC Example

Recognition © Kenneth Dawson-Howe 2015 Slide 44

44

RANSAC Examples

From <http://web.engr.illinois.edu/~nanchen2/cs498dwh/proj/>

Recognition Slide 46

46

Features

- Introduction
- Moravec Corner Detection
- Harris/Plessey Corner Detection
- Scale Invariant Feature Transform (SIFT)

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Features Slide 1

1

Introduction – Possible interpretations

The Aperture Problem.

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Features Slide 3

3

Introduction – Where have the edges gone?

- Given two images (a) and (b) taken at different times determine the movement of edge points from frame to frame...

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Features Slide 2

2

Introduction – Using Features / Corners instead

- Use corners / image features / interest points
- Corner = intersection of two edges
- Interest point = any feature which can be robustly detected

- Reduces number of points
- Easier to establish correspondences
- Spurious features

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Features Slide 4

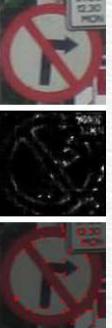
4

1

2

Introduction – Steps for corner detection

1. Determine cornerness values.
 - For each pixel
 - Main difference
 - Produces a Cornerness map.
2. Non-maxima suppression.
 - Multiple responses
 - Compare to local neighbours
3. Threshold the cornerness map.
 - Significant corners.

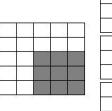


Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Features Slide 5

5

Moravec – Binary Example

Corner:	Edge:
	
Minimum difference: 2	Minimum difference: 0

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Features Slide 7

7

Moravec corner detection



- Looks at the local variation around a point
- Compares local image patches

$$V_{u,v}(i,j) = \sum_{\forall a,b \in Window} (f(i+u+a, j+v+b) - f(i+a, j+b))^2$$

where $(u,v) \in \{(-1,-1), (-1,0), (-1,1), (0,-1), (0,1), (1,-1), (1,0), (1,1)\}$ and the Window is typically 3x3, 5x5 or 7x7

- Select the Minimum value of $V_{u,v}(i,j)$

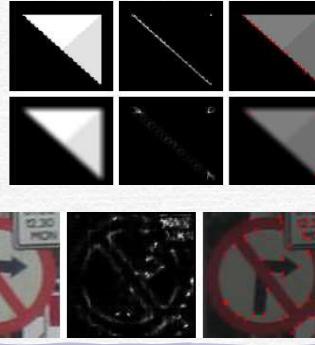
Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Features Slide 6

6

Moravec – Flaws.

- Anisotropic response
- Diagonal lines
- Smoothing



- Noisy response
- Larger area
- Smoothing

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Features Slide 8

8

Harris / Plessey corner detection

- Difference – Cornerness determination
- Uses
 - Partial derivatives
 - Gaussian weighting
 - Matrix Eigenvalues

```
Ptr<FeatureDetector> harris_detector =
    GFTTDetector::create( 1000, 0.01, 10, 3, true );
vector<KeyPoint> keypoints;
harris_detector->detect( gray_image, keypoints );
```

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Features Slide 9

9

Harris / Plessey corner detection

- From the matrix we can compute the Eigenvalues
 - Both high => corner
 - One high => edge
 - None high => constant region
- Harris & Stephens proposed the following cornerness measure:

$$M = \begin{bmatrix} \sum_{(i,j) \in W} \left(\frac{\delta f(i,j)}{\delta i} \right)^2 & \sum_{(i,j) \in W} \frac{\delta f(i,j) \delta f(i,j)}{\delta i \delta j} \\ \sum_{(i,j) \in W} \frac{\delta f(i,j) \delta f(i,j)}{\delta i \delta j} & \sum_{(i,j) \in W} \left(\frac{\delta f(i,j)}{\delta j} \right)^2 \end{bmatrix} = \begin{bmatrix} A & B \\ B & C \end{bmatrix}$$

0.04	0.12	0.04
0.12	0.36	0.12
0.04	0.12	0.04

$\det(M) = \lambda_1 \lambda_2 = AC + B^2$
 $\text{trace}(M) = \lambda_1 + \lambda_2 = A + C$
 $C(i,j) = \det(M) - k(\text{trace}(M))^2$

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Features Slide 11

11

Harris / Plessey corner detection

- Consider the intensity variation for an arbitrary shift $(\Delta i, \Delta j)$ as

$$SSD_W(\Delta i, \Delta j) = \sum_{(i,j) \in W} (f(i,j) - f(i - \Delta i, j - \Delta j))^2$$

- If $f(i - \Delta i, j - \Delta j) \approx f(i,j) + \left[\frac{\delta f(i,j)}{\delta i} \frac{\delta f(i,j)}{\delta j} \right] [\Delta i \Delta j]$
- Then $SSD_W(\Delta i, \Delta j) = \sum_{(i,j) \in W} \left(f(i,j) - f(i,j) - \left[\frac{\delta f(i,j)}{\delta i} \frac{\delta f(i,j)}{\delta j} \right] [\Delta i \Delta j] \right)^2$

$$SSD_W(\Delta i, \Delta j) = \sum_{(i,j) \in W} \left(\left[\frac{\delta f(i,j)}{\delta i} \frac{\delta f(i,j)}{\delta j} \right] [\Delta i \Delta j] \right)^2$$

$$SSD_W(\Delta i, \Delta j) = \sum_{(i,j) \in W} \left([\Delta i \Delta j] \left[\begin{bmatrix} \frac{\delta f(i,j)}{\delta i} & \frac{\delta f(i,j)}{\delta j} \\ \frac{\delta f(i,j)}{\delta i} & \frac{\delta f(i,j)}{\delta j} \end{bmatrix} \right] [\Delta i \Delta j] \right)$$

$$SSD_W(\Delta i, \Delta j) = [\Delta i \Delta j] \begin{bmatrix} \sum_{(i,j) \in W} \left(\frac{\delta f(i,j)}{\delta i} \right)^2 & \sum_{(i,j) \in W} \frac{\delta f(i,j) \delta f(i,j)}{\delta i \delta j} \\ \sum_{(i,j) \in W} \frac{\delta f(i,j) \delta f(i,j)}{\delta i \delta j} & \sum_{(i,j) \in W} \left(\frac{\delta f(i,j)}{\delta j} \right)^2 \end{bmatrix} [\Delta i \Delta j]$$

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Features Slide 10

10

Harris / Plessey corner detection

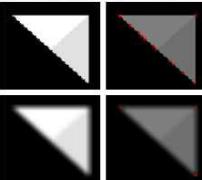
Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Features Slide 12

12

Harris / Plessey – Pros and Cons

- Cons:
 - More expensive computationally
 - Sensitive to noise
 - Somewhat anisotropic
- Pros:
 - Very repeatable response
 - Better detection rate



```
Mat display_image;
drawKeypoints( image, keypoints, display_image,
               Scalar( 0, 0, 255 ) );
```

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Features Slide 13

13

Scale Invariant Feature Transform (SIFT)

- FROM "Distinctive Image Features from Scale-Invariant Keypoints",
by David G. Lowe in International Journal of Computer Vision, 60, 2 (2004), pp.91-110
- Motivation:
 - Providing repeatable robust features for
 - Tracking,
 - Recognition,
 - Panorama Stitching, etc.
- Features:
 - Invariant to scaling, & rotation.
 - Partly invariant to illumination and viewpoint changes

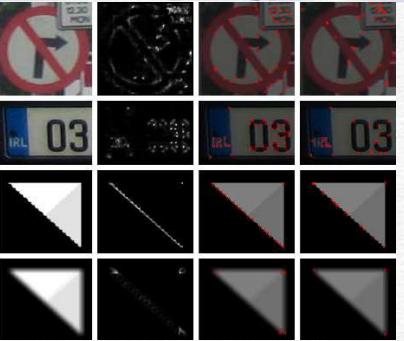


Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Features Slide 15

15

Moravec & Harris/Plessey



Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Features Slide 14

14

SIFT – Overview & Contents

1. Scale Space Extrema Detection
 - Scale Space
 - Difference of Gaussian
 - Locate Extrema
2. Accurate Keypoint Location
 - Sub-pixel locate
 - Filter response – remove low contrast and features primarily along an edge
3. Keypoint Orientation assignment
4. Keypoint Descriptors
- Matching Descriptors – including dropping poor ones
- Applications

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Features Slide 16

16

SIFT – OpenCV code

```
// The following code is valid up to v2.4.11:
Ptr<FeatureDetector> feature_detector =
    FeatureDetector::create("SIFT");
vector<KeyPoint> keypoints;
feature_detector->detect( gray_image, keypoints );
```

Features Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014 Slide 17

17

SIFT – Scale Space Extrema Detection

- Stable keypoint locations defined to be at extrema in the Difference of Gaussian (DoG) functions across scale space...
- $D(x,y,\sigma) = L(x,y,k\sigma) - L(x,y,\sigma)$

Extrema...

- Centre point is Min or Max of
- Local 3x3 region in current DoG
- and in adjacent scales
- IN any octave

Features Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014 Slide 19

19

SIFT – Scale Space Extrema Detection

- For scale invariance, consider the image at multiple scales
- $L(x,y,\sigma) = G(x,y,\sigma) * I(x,y)$
- Applied in different octaves of scale space
- Each octave corresponds to a doubling of σ

Features Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014 Slide 18

18

SIFT – Accurate Keypoint Location

- Originally location and scale taken from central point
- Locate keypoints more precisely
 - Model data locally using a 3D quadratic
 - Locate interpolated maximum/minimum

Features Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014 Slide 20

20

SIFT – Accurate Keypoint Location

- Discard low contrast keypoints
 - If the local contrast is too low discard the keypoint
 - Evaluated from the curvature of the 3D quadratic...

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Features Slide 21

21

SIFT – Keypoint Orientation

- For scale invariance, the keypoint scale is used to select the smoothed image with the closest scale
- For orientation invariance we describe the keypoint wrt. the principal orientation
 - Create an orientation histogram (36 bins)
 - $m(x, y) = \sqrt{(L(x+1, y) - L(x-1, y))^2 + (L(x, y+1) - L(x, y-1))^2}$
 - $\theta(x, y) = \tan^{-1}((L(x, y+1) - L(x, y-1))/(L(x+1, y) - L(x-1, y)))$
 - Weight by gradient magnitude
 - Sample points around the keypoint
 - Highest peak + peaks within 80%
 - Oriented keypoint(s)
 - Stable results....

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Features Slide 23

23

SIFT – Accurate Keypoint Location

- Discard poorly localised keypoints (e.g. along an edge)

$$\mathbf{H} = \begin{bmatrix} D_{xx} & D_{xy} \\ D_{xy} & D_{yy} \end{bmatrix} \quad \text{Tr}(\mathbf{H}) = D_{xx} + D_{yy} = \alpha + \beta$$

$$\text{Det}(\mathbf{H}) = D_{xx}D_{yy} - (D_{xy})^2 = \alpha\beta \quad \alpha = r\beta$$

$$\frac{\text{Tr}(\mathbf{H})^2}{\text{Det}(\mathbf{H})} = \frac{(\alpha + \beta)^2}{\alpha\beta} = \frac{(r\beta + \beta)^2}{r\beta^2} = \frac{(r+1)^2}{r}$$

$$\frac{\text{Tr}(\mathbf{H})^2}{\text{Det}(\mathbf{H})} < \frac{(r+1)^2}{r}$$

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Features Slide 22

22

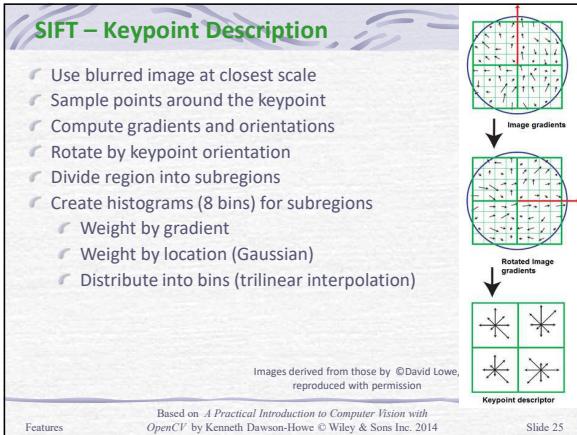
SIFT – Keypoint Description

- Could sample image intensity at relevant scale
- Match using normalized cross correlation
- Sensitive to
 - affine transformations,
 - 3D viewpoint changes and
 - non-rigid deformations
- A better approach (Edelman et al. 1997)
 - Based on a model of biological vision
 - Consider gradients at particular orientations and spatial frequencies
 - Location not required to be precise

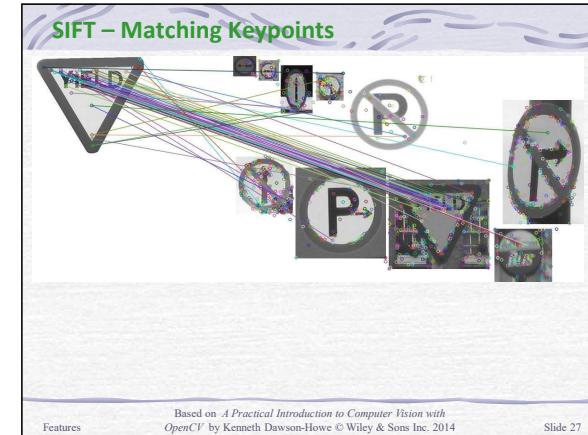
Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Features Slide 24

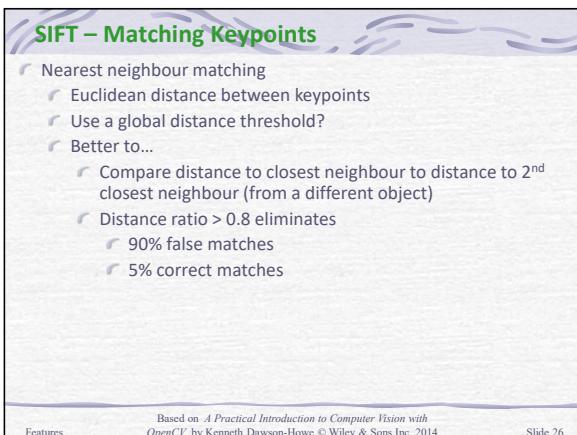
24



25



27



26

SIFT – OpenCV code

```
// The following code is valid up to v2.4.11:
SiftFeatureDetector sift_detector;
vector<KeyPoint> keypoints1, keypoints2;
sift_detector.detect( gray_image1, keypoints1 );
sift_detector.detect( gray_image2, keypoints2 );
// Extract feature descriptors
SiftDescriptorExtractor sift_extractor;
Mat descriptors1, descriptors2;
sift_extractor.compute( gray_image1, keypoints1, descriptors1 );
sift_extractor.compute( gray_image2, keypoints2, descriptors2 );
...
```

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Features Slide 28

28

SIFT – OpenCV code



```

...
// Match descriptors.
BFMatcher sift_matcher(NORM_L2);
vector< DMatch > matches;
matcher.match( descriptors1, descriptors2, matches );
// Display SIFT matches
Mat display_image;
drawMatches( gray_image1, keypoints1, gray_image2,
keypoints2, matches, display_image );

```

Features Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014 Slide 29

29

SIFT – Application: Recognition



Images ©David Lowe, reproduced with permission

Features Based on *A Practical Introduction to Computer Vision with OpenCV 3* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014 Slide 31

31

SIFT – Recognition

- ➊ Match at least 3 features
- ➋ Cluster matches
 - ➌ Hough transform
 - ➍ Location (2D)
 - ➎ Scale
 - ➏ Orientation
 - ➐ Really a 6D problem
 - ➑ Use broad bins
 - ➒ 30° for orientation
 - ➓ Factor of 2 for scale
 - ➔ 0.25 times image dimension for location
 - ➑ Consider all bins with at least 3 entries
 - ➒ Determine affine transformation

Features Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014 Slide 30

30

SIFT – Application: Recognition



Images ©David Lowe, reproduced with permission

Features Based on *A Practical Introduction to Computer Vision with OpenCV 3* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014 Slide 32

32

15

16

SIFT – Application: Panorama Stitching

Features

Images © Lucas Mach, released under CC-BY 3.0

Based on *A Practical Introduction to Computer Vision with OpenCV 3* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 33

33

Geometric

- Geometric transformations
- Pixel co-ordinate transformationns
- Brightness interpolation

Geometric

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 1

1

SIFT – Application: Tracking

Features

Based on *A Practical Introduction to Computer Vision with OpenCV 3* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 34

34

Geometric transformations

- Computer graphics
 - Introduce distortion
- Image processing
 - Image mosaicing
 - Matching/Registering image
 - Eliminating distortion
 - Simplifying further processing
 - e.g. OCR

[Examples](#)

Geometric

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 2

2

Geometric transformations - Specification

- Problem:
 - Distorted image $f(i, j)$
 - Corrected image $f'(i', j')$
- Mapping: $i = T_i(i', j')$ $j = T_j(i', j')$

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Geometric

Slide 3

3

Output → Input

- A backwards transformation?
 - Why aren't we mapping from the input image to the output?
 - For each point maintain the new continuous coordinates
 - For each output point
 - Find the nearest transformed point(s)
 - Interpolate a value from this point(s)
- We would have to
 - Store all real coordinates
 - Search for nearest values for each output pixel
 - Interpolate in corrected image space

Geometric

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 5

5

Geometric transformations - Specification

- Application steps:
 - Define the transformation
 - Known in advance
 - Determine through correspondences
 - Image to known
 - Image to image
 - Apply the transformation
 - For every point in the output image
 - Determine where it came from using T
 - Interpolate a value for the output point

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Geometric

Slide 4

4

Affine transformations

- Definition
$$\begin{bmatrix} i \\ j \end{bmatrix} = \begin{bmatrix} a_{00} & a_{01} & a_{02} \\ a_{10} & a_{11} & a_{12} \end{bmatrix} \begin{bmatrix} i' \\ j' \\ 1 \end{bmatrix}$$
- Known transformations...
 - E.g. Translation
$$\begin{bmatrix} i \\ j \end{bmatrix} = \begin{bmatrix} 1 & 0 & m \\ 0 & 1 & n \end{bmatrix} \begin{bmatrix} i' \\ j' \\ 1 \end{bmatrix}$$
- Unknown transformations...
 - Requires at least 3 observations

```
Mat affine_matrix( 2, 3, CV_32FC1 );
...
warpAffine( image, result, affine_matrix, image.size() );
```

Geometric

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 6

6

Simple affine transformations

- Rotation

$$\begin{bmatrix} i \\ j \end{bmatrix} = \begin{bmatrix} \cos \phi & \sin \phi & 0 \\ -\sin \phi & \cos \phi & 0 \end{bmatrix} \begin{bmatrix} i' \\ j' \\ 1 \end{bmatrix}$$
- Change of scale

$$\begin{bmatrix} i \\ j \end{bmatrix} = \begin{bmatrix} a & 0 & 0 \\ 0 & b & 0 \end{bmatrix} \begin{bmatrix} i' \\ j' \\ 1 \end{bmatrix}$$
- Skewing

$$\begin{bmatrix} i \\ j \end{bmatrix} = \begin{bmatrix} 1 & \tan \phi & 0 \\ 0 & 1 & 0 \end{bmatrix} \begin{bmatrix} i' \\ j' \\ 1 \end{bmatrix}$$

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Geometric Slide 7

7

Unknown affine transformations

- Given...

$$\begin{bmatrix} i \\ j \end{bmatrix} = \begin{bmatrix} a_{00} & a_{01} & a_{02} \\ a_{10} & a_{11} & a_{12} \end{bmatrix} \begin{bmatrix} i' \\ j' \\ 1 \end{bmatrix}$$
- If we have 3 observations...

$$(i_1, j_1) \leftrightarrow (i'_1, j'_1)$$

$$(i_2, j_2) \leftrightarrow (i'_2, j'_2)$$

$$(i_3, j_3) \leftrightarrow (i'_3, j'_3)$$
- We can reorganise...

$$\begin{bmatrix} i_1 \\ j_1 \\ 1 \end{bmatrix} = \begin{bmatrix} i'_1 & j'_1 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & i'_1 & j'_1 & 1 \\ i'_2 & j'_2 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & i'_2 & j'_2 & 1 \\ i'_3 & j'_3 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & i'_3 & j'_3 & 1 \end{bmatrix} \begin{bmatrix} a_{00} \\ a_{01} \\ a_{02} \\ a_{10} \\ a_{11} \\ a_{12} \end{bmatrix}$$
- And take the inverse to compute the a coefficients

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Geometric Slide 9

9

More affine transformations

- Panoramic distortion

$$\begin{bmatrix} i \\ j \end{bmatrix} = \begin{bmatrix} a & 0 & 0 \\ 0 & 1 & 0 \end{bmatrix} \begin{bmatrix} i' \\ j' \\ 1 \end{bmatrix}$$

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Geometric Slide 8

8

Unknown affine transformations

Point2f source [3], destination [3];
...
`affine_matrix = getAffineTransform(source,destination);`

- More observations
 - Better estimate of the a coefficients
 - Must use the psuedo inverse

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Geometric Slide 10

10

Perspective transformations

↳ Perspective projection

- Planar surface
- Not parallel to the image plane.
 - Cannot be corrected with the affine transformation
- Need a perspective transformation

$$\begin{bmatrix} i.w \\ j.w \\ w \end{bmatrix} = \begin{bmatrix} p_{00} & p_{01} & p_{02} \\ p_{10} & p_{11} & p_{12} \\ p_{20} & p_{21} & 1 \end{bmatrix} \begin{bmatrix} i' \\ j' \\ 1 \end{bmatrix}$$

Geometric Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014 Slide 11

11

Perspective transformations

Then

$$\begin{bmatrix} i_1 \\ j_1 \\ i_2 \\ j_2 \\ i_3 \\ j_3 \\ i_4 \\ j_4 \end{bmatrix} = \begin{bmatrix} i'_1 & j'_1 & 1 & 0 & 0 & 0 & -i_1i'_1 & -i_1j'_1 \\ 0 & 0 & 0 & i'_1 & j'_1 & 1 & -j_1i'_1 & -j_1j'_1 \\ i'_2 & j'_2 & 1 & 0 & 0 & 0 & -i_2i'_2 & -i_2j'_2 \\ 0 & 0 & 0 & i'_2 & j'_2 & 1 & -j_2i'_2 & -j_2j'_2 \\ i'_3 & j'_3 & 1 & 0 & 0 & 0 & -i_3i'_3 & -i_3j'_3 \\ 0 & 0 & 0 & i'_3 & j'_3 & 1 & -j_3i'_3 & -j_3j'_3 \\ i'_4 & j'_4 & 1 & 0 & 0 & 0 & -i_4i'_4 & -i_4j'_4 \\ 0 & 0 & 0 & i'_4 & j'_4 & 1 & -j_4i'_4 & -j_4j'_4 \end{bmatrix} \begin{bmatrix} p_{00} \\ p_{01} \\ p_{02} \\ p_{10} \\ p_{11} \\ p_{12} \\ p_{20} \\ p_{21} \end{bmatrix}$$

↳ Multiply by the inverse of the square matrix...

↳ More observations gives

- A more accurate transformation
- Matrix becomes non-square...
- Pseudo inverse required

Geometric Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014 Slide 13

13

Perspective transformations

From

$$\begin{bmatrix} i.w \\ j.w \\ w \end{bmatrix} = \begin{bmatrix} p_{00} & p_{01} & p_{02} \\ p_{10} & p_{11} & p_{12} \\ p_{20} & p_{21} & 1 \end{bmatrix} \begin{bmatrix} i' \\ j' \\ 1 \end{bmatrix}$$

↳ We know $i.w = p_{00}.i' + p_{01}.j' + p_{02}$
 $w = p_{20}.i' + p_{21}.j' + 1$

↳ Hence $i = p_{00}.i' + p_{01}.j' + p_{02} - p_{20}.i.i' - p_{21}.i.j'$

↳ And similarly $j = p_{10}.i' + p_{11}.j' + p_{12} - p_{20}.j.i' - p_{21}.j.j'$

↳ If we observe 4 mappings...

Geometric Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014 Slide 12

12

Perspective transformations

Point2f source [4], destination [4];
// Assign values to source and destination points.
perspective_matrix = getPerspectiveTransform(source, destination);
warpPerspective(image, result, perspective_matrix, result.size());

Geometric Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014 Slide 14

14

More complex transformations

- Approximation by a polynomial

$$i = T_i(i', j') = a_{00} + a_{10}i' + a_{01}j' + a_{11}i'j' + a_{02}(j')^2 + a_{20}(i')^2 + a_{12}i'(j')^2 + a_{21}(i')^2j' + a_{22}(i')^2(j')^2 + \dots$$

$$j = T_j(i', j') = b_{00} + b_{10}i' + b_{01}j' + b_{11}i'j' + b_{02}(j')^2 + b_{20}(i')^2 + b_{12}i'(j')^2 + b_{21}(i')^2j' + b_{22}(i')^2(j')^2 + \dots$$
 - No. of correspondences.
 - Distribution of points
 - Solve simultaneous linear equations
 - Least squared error solution
- Even more complex
 - Partition the image.

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Geometric

Slide 15

15

Nearest neighborhood interpolation

- Simple formulation: $f'(i', j') = f(\text{round}(i), \text{round}(j))$
- Error is perceptible
 - Step edges

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Geometric

Slide 17

17

Brightness interpolation

- Locations are not integer coordinates
- Interpolate output pixel value from
 - The nearby pixels in the original (uncorrected) image
- Loss of accuracy
 - Complexity of interpolation scheme
- Interpolation schemes:
 - Nearest neighbour
 - Bilinear interpolation
 - Bicubic interpolation

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Geometric

Slide 16

16

Bi-linear interpolation

$$f'(i', j') = (\text{trunc}(i) + 1 - i)(\text{trunc}(j) + 1 - j)f(\text{trunc}(i), \text{trunc}(j)) + (i - \text{trunc}(i))(\text{trunc}(j) + 1 - j)f(\text{trunc}(i) + 1, \text{trunc}(j)) + (\text{trunc}(i) + 1 - i)(j - \text{trunc}(j))f(\text{trunc}(i), \text{trunc}(j) + 1) + (i - \text{trunc}(i))(j - \text{trunc}(j))f(\text{trunc}(i) + 1, \text{trunc}(j) + 1)$$

- Brightness function is bilinear
- Neighbours contribute
- Error is perceptible
 - Blurring

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Geometric

Slide 18

18

Bi-cubic interpolation

- ✓ Approximate using a bi-cubic polynomial surface
- ✓ No step-like boundary problems
- ✓ Copes with linear interpolation blurring
- ✓ Often used in raster displays...



Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Geometric

Slide 19

19

Camera models – Distortion

- ✓ Radial distortion
 - Radially symmetric from the optical axis
 - Change in magnification
 - Barrel distortion
 - Pincushion distortion
 - Caused by the lenses
$$i' = i(1 + k_1r^2 + k_2r^4 + k_3r^6)$$

$$j' = j(1 + k_1r^2 + k_2r^4 + k_3r^6)$$

$$r = \sqrt{i^2 + j^2}$$
- ✓ Tangential distortion
 - Uneven magnification from one side to the other
 - Lenses not parallel to the image plane
$$i' = i + (2p_1ij + p_2(r^2 + 2i^2))$$

$$j' = j + (2p_2ij + p_1(r^2 + 2j^2))$$

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Geometric

Slide 21

21

Brightness interpolation

```
int interpolation_scheme = INTER_LINEAR;
// Nearest neighbour interpolation: INTER_NEAREST
// Bilinear interpolation: INTER_LINEAR
// Bicubic interpolation: INTER_CUBIC
warpAffine( image, result, affine_matrix, result_size,
            interpolation_scheme );
warpPerspective( image, result, perspective_matrix,
                  result_size, interpolation_scheme );
```

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

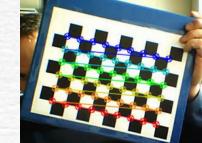
Geometric

Slide 20

20

Camera models – Removing distortion

- ✓ Calibrate using multiple images
 - Known object
 - Different positions
 - Computes camera matrix
 - & distortion parameters
- ✓ Remove distortion
 - Mainly using the distortion parameters



```
calibrateCamera( object_points, image_points,
                  image_size, camera_matrix, distortion_coefficients,
                  rotation_vectors, translation_vectors );
...
undistort( camera_image, corrected_image,
           camera_matrix, distortion_coefficients );
```

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Geometric

Slide 22

10

11

Histograms

- 1D Histograms
- 3D Histograms
- Equalisation
- Histogram Comparison
- Back Projection

Histograms Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014 Slide 1

1D Histograms

Is this useful?

- Global information
- Useful for classification ?
- Not unique

Histograms Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014 Slide 3

1D Histograms

Determine the frequency of brightness values

- Initialise:
 - $h(z) = 0$ for all values of z
- Compute:
 - For all pixels (i,j) : $h(f(i,j))++$

Histograms Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014 Slide 2

1D Histograms

```

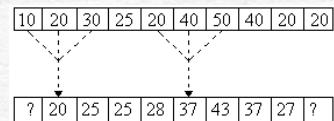
Mat display_image;
MatND gray_histogram;
const int* channel_numbers = { 0 };
float channel_range[] = { 0.0, 255.0 };
const float* channel_ranges = channel_range;
int number_bins = 64;
calcHist( &gray_image, 1, channel_numbers, Mat(),
          gray_histogram, 1, &number_bins, &channel_ranges );
OneDHistogram::Draw1DHistogram( &gray_histogram, 1,
                                display_image );
imshow("Greyscale histogram", display_image);

```

Histograms Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014 Slide 4

1D Histograms – Smoothing

- Local minima and maxima are useful.
- But how can we deal with noise though?



- Smooth the histogram...
 - For all values v : $h_{\text{new}}(v) = (h(v-1) + h(v) + h(v+1)) / 3$
- What do we do at the ends?
 - Do not compute OR Wraparound OR Duplicate values
OR Reflect values OR Constant values ??

Histograms

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 5

1D Histograms – Colour Histograms

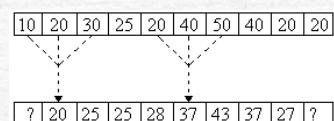
- Determine histograms for each channel independently...
- Choice of colour space...

Histograms

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 7

1D Histograms – Smoothing



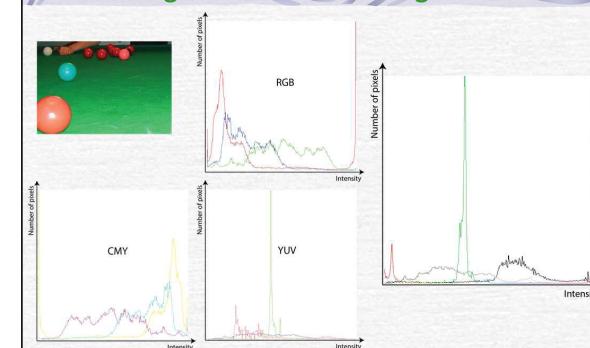
```
MatND smoothed_histogram = histogram[channel].clone();
for(int i = 1; i < histogram[channel].rows - 1; ++i)
    smoothed_histogram[channel].at<float>(i) =
        (histogram.at<float>(i-1) + histogram.at<float>(i) +
         histogram.at<float>(i+1)) / 3;
```

Histograms

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 6

1D Histograms – Colour Histograms



Histograms

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 8

1D Histograms – Colour Histograms

```
MatND* histogram = new MatND[image.channels()];
vector<Mat> channels(image.channels());
split( image, channels );
const int* channel_numbers = { 0 };
float channel_range[] = { 0.0, 255.0 };
const float* channel_ranges = channel_range;
int number_bins = 64;
for (int chan=0; chan < image.channels(); chan++)
    calcHist( &(channels[chan]), 1, channel_numbers, Mat(),
              histogram[chan], 1, &number_bins, &channel_ranges );
OneDHistogram::Draw1DHistogram(histogram,
                                image.channels(), display_image );
```

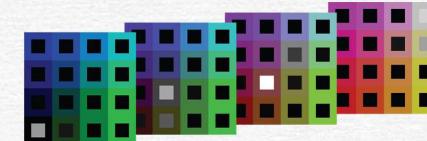
Histograms

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 9

3D Histograms

- Reduce quantisation
 - 6 bits = 262,144
 - 4 bits = 4,096
 - 2 bits = 64



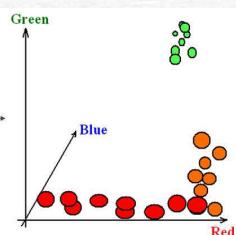
Histograms

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 11

3D Histograms

- Channels are not independent
- Better discrimination comes from considering all channels simultaneously
- Number of cells?
 - 8 bits = 16,777,216



Histograms

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 10

3D Histograms

```
MatND histogram;
int channel_numbers[] = { 0, 1, 2 };
int* number_bins = new int[image.channels()];
for (ch=0; ch < image.channels(); ch++)
    number_bins[ch] = 64;
float ch_range[] = { 0.0, 255.0 };
const float* channel_ranges[] = {ch_range, ch_range, ch_range};
calcHist( &image, 1, channel_numbers, Mat(), histogram,
          image.channels(), a_number_bins, channel_ranges );
```

Histograms

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 12

Equalisation

- If an image has insufficient contrast
- Human can distinguish 700-900 greyscales
- Evenly distribute the greyscales...
 - Result has missing greyscales
- Normally equalise only the greyscales / luminance

Histograms Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014 Slide 13

Equalisation

```
vector<Mat> channels( hls_image.channels() );
split(hls_image, channels);
equalizeHist( channels[1], channels[1] );
merge( channels, hls_image );
```

Histograms Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014 Slide 15

Equalisation

```
// Create a lookup table to map the luminances
// h[x] = histogram of luminance values image f(i,j).
pixels_so_far = 0
num_pixels = image.rows * image.cols
output = 0
for input = 0 to 255
  pixels_so_far = pixels_so_far + h[ input ]
  new_output = (pixels_so_far*256) / (num_pixels+1)
  LUT[ input ] = (output+1+new_output) / 2
  output = new_output
// Apply the lookup table LUT(x) to the image:
for every pixel f(i,j)
  f'(i,j) = LUT[ f(i,j) ]
```

Histograms Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014 Slide 14

Histogram Comparison

- To find similar images...
 - Use metadata
 - Compare images
- Compare the colour distributions
- Need a metric for comparisons...

Histograms Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014 Slide 16



Histogram Comparison – Earth Mover

- An alternative to the metrics is to compute the Earth Mover's Distance...
 - Minimum cost for turning a distribution into another distribution
 - Compare images
- 1D solution:
 - $\text{EMD}(-1) = 0$
 - $\text{EMD}(i) = h_1(i) + \text{EMD}(i-1) - h_2(i)$
 - Earth Mover's Distance = $\sum_i |\text{EMD}(i)|$
- Colour EMD is harder to compute...

Histograms Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014 Slide 19

Histogram Comparison – Metrics

- $D_{Correlation}(h_1, h_2) = \frac{\sum_i (h_1(i) - \bar{h}_1)(h_2(i) - \bar{h}_2)}{\sqrt{\sum_i (h_1(i) - \bar{h}_1)^2 \sum_i (h_2(i) - \bar{h}_2)^2}}$
- $D_{Chi-Square}(h_1, h_2) = \sum_i \frac{(h_1(i) - h_2(i))^2}{(h_1(i) + h_2(i))}$
- $D_{Intersection}(h_1, h_2) = \sum_i \min(h_1(i), h_2(i))$
- $D_{Bhattacharyya}(h_1, h_2) = \sqrt{1 - \frac{1}{\sqrt{h_1 \cdot h_2 \cdot N^2}} \sum_i \sqrt{h_1(i) \cdot h_2(i)}}$
- where
 - N is the number of bins in the histograms,
 - $\bar{h}_k = \sum_i (h_k(i)) / N$

Histograms Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014 Slide 18

Histogram Comparison

```
normalize( histogram1, histogram1, 1.0);
normalize( histogram2, histogram2, 1.0);
double matching_score = compareHist( histogram1,
                                      histogram2, CV_COMP_CORREL);

We can also use Chi-Square (CV_COMP_CHISQR),
Intersection (CV_COMP_INTERSECT) or
Bhattacharyya (CV_COMP_BHATTACHARYYA) metrics or
alternatively can use the Earth Mover's Distance function
(EMD())
```

Histograms Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014 Slide 20

Histogram Back Projection

- A better approach to selecting colours (based on samples):
 1. Obtain a representative sample set of the colours.
 2. Histogram those samples.
 3. Normalize that histogram so that the maximum value is 1.0.
 4. Back project the normalized histogram onto any image $f(i,j)$.
 5. This will provide a 'probability' image $p(i,j)$ which indicates the similarity between $f(i,j)$ and the sample set.



Histograms

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 21

Images

- Camera models
- Digital images
- Colour images
- Noise
- Smoothing

Images

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 1

1

Histogram Back Projection



```
calcHist( &hls_samples_image, 1, channel_numbers, Mat(),
          histogram,image.channels(),number_bins,channel_ranges);
normalize( histogram, histogram, 1.0);
Mat probabilities = histogram.BackProject( hls_image );
```

Histograms

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

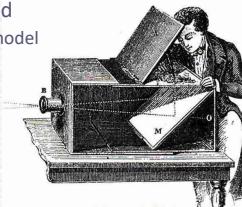
Slide 22

Camera models

- Components:
 - A photosensitive image plane
 - A housing
 - A lenses

- Mathematical model needed

- The simple pinhole camera model
- Distortions

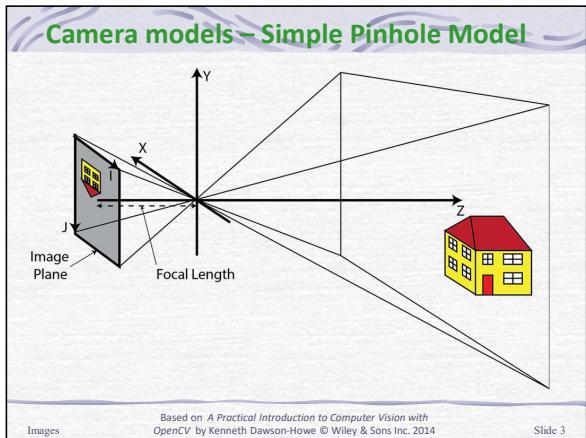


Images

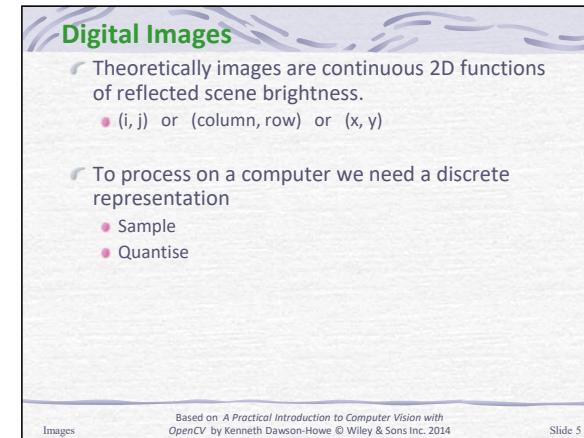
Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 2

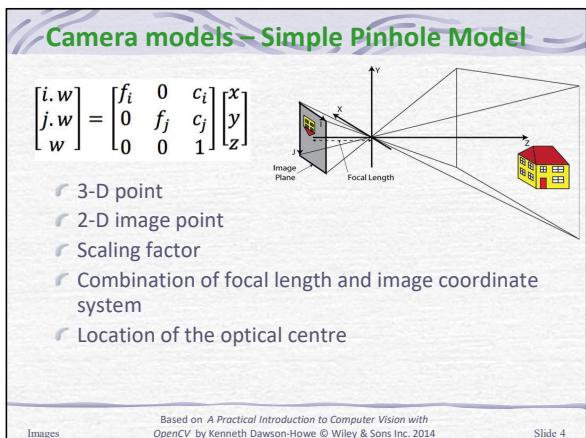
2



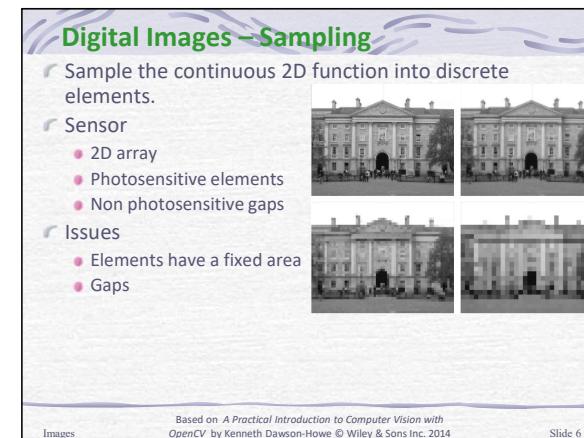
3



5



4



6

Digital Images – Sampling

- How many samples do we need ?

- Wasted space and computation time
- Enough for the objects of interest



```
Mat image, smaller_image;  
resize( image, smaller_image,  
Size( image1.cols/2, image.rows/2 ));
```

Images

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 7

Digital Images – Quantisation

- How many bits do we need?

- Wasted space ?
- Losing the ability to distinguish objects



```
void ChangeQuantisationGrey( Mat &image, int num_bits )  
{  
    CV_Assert( (image.type() == CV_8UC1) && (num_bits >= 1) &&  
              (num_bits <= 8) );  
    uchar mask = 0xFF << (8-num_bits);  
    for (int row=0; row < image.rows; row++)  
        for (int col=0; col < image.cols; col++)  
            image.at<uchar>(row,col) = image.at<uchar>(row,col) & mask;  
}
```

Images

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 9

9

Digital Images – Quantisation

- Represent the individual image points as digital values.

- Typically 8 bits



Images

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 8

Colour Images

- Luminance only
 - Simple representation
 - Humans can understand
- Colour images (luminance + chrominance)
 - Multiple channels (typically 3)
 - Around 16.8 million colours
 - More complex to process
 - Facilitate certain operations



Images

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 10

10

Colour Images – Processing

```
void InvertColour( Mat& input_image, Mat& output_image )
{
    CV_Assert( input_image.type() == CV_8UC3 );
    output_image = input_image.clone();
    for (int row=0; row < input_image.rows; row++)
        for (int col=0; col < input_image.cols; col++)
            for (int channel=0; channel <
                input_image.channels(); channel++)
                output_image.at<Vec3b>(row,col)[channel] = 255 -
                input_image.at<Vec3b>(row,col)[channel];
}
```

Images

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

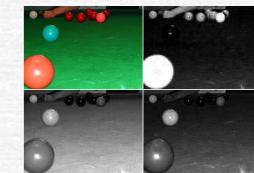
Slide 11

11

Colour Images – RGB Images

Red-Green-Blue images

- Most common
- Channels correspond roughly to
 - Red (700nm)
 - Green (546nm)
 - Blue (436nm)
- Colours combined in display



Images

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 13

13

Colour Images – Efficient processing

```
int image_rows = image.rows;
int image_columns = image.cols;
for (int row=0; row < image_rows; row++) {
    uchar* value = image.ptr<uchar>(row);
    uchar* result_value = result_image.ptr<uchar>(row);
    for (int column=0; column < image_columns; column++)
    {
        *result_value++ = *value++ ^ 0xFF;
        *result_value++ = *value++ ^ 0xFF;
        *result_value++ = *value++ ^ 0xFF;
    }
}
```

Images

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 12

12

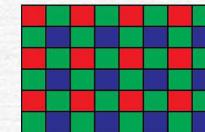
Colour Images – RGB Images

Converting to Greyscale

- $Y = 0.299R + 0.587G + 0.114B$

Camera photosensitive elements

- Separate Red, Green & Blue elements
- Sometimes sensitive to all visible wavelengths
- Bayer pattern:



Images

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 14

14

Colour Images – RGB Images

```
Mat bgr_image, grey_image;
cvtColor(bgr_image, grey_image, CV_BGR2GRAY);
vector<Mat> bgr_images(3);
split(bgr_image, bgr_images);
Mat& blue_image = bgr_images[0];
```

Images

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 15

15

Colour Images – YUV Images

- Used for analogue television signals

- PAL, NTSC
- 4 Y to 1 U to 1 V

- Conversion from RGB

$$\begin{aligned}Y &= 0.299R + 0.587G + 0.114B \\U &= 0.492 * (B-Y) \\V &= 0.877 * (R-Y)\end{aligned}$$



Images

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

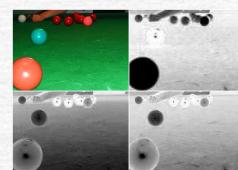
Slide 17

17

Colour Images – CMY Images

- Cyan-Magenta-Yellow images

- Secondary colours
- Subtractive colour scheme
 - $C = 255 - R$
 - $M = 255 - G$
 - $Y = 255 - B$
- Often used in printers



CMY is not directly supported in OpenCV.

Images

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

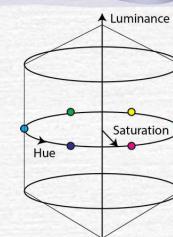
Slide 16

16

Colour Images – HLS Images

- Hue-Luminance-Saturation images

- Separates Luminance & Chrominance
- Values we humans can relate to...
 - Hue 0°..360°
 - Luminance 0..1
 - Saturation 0..1
- Watch out for circular Hue...



Images

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 18

18

Colour Images – HLS Images

- Conversion from RGB

$$L = \frac{\text{Max}(R, G, B) + \text{Min}(R, G, B)}{2}$$

$$S = \begin{cases} \frac{\text{Max}(R, G, B) - \text{Min}(R, G, B)}{\text{Max}(R, G, B) + \text{Min}(R, G, B)} & \text{if } L < 0.5 \\ \frac{\text{Max}(R, G, B) - \text{Min}(R, G, B)}{2 - (\text{Max}(R, G, B) + \text{Min}(R, G, B))} & \text{if } L \geq 0.5 \end{cases}$$

$$H = \begin{cases} \frac{60 \cdot (G - B)}{S} & \text{if } R = \text{Max}(R, G, B) \\ \frac{120 + 60 \cdot (B - R)}{S} & \text{if } G = \text{Max}(R, G, B) \\ \frac{240 + 60 \cdot (R - G)}{S} & \text{if } B = \text{Max}(R, G, B) \end{cases}$$

```
cvtColor(bgr_image, hls_image, CV_BGR2HLS);
// Hue ranges from 0 to 179.
```

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Images Slide 19

19

Noise

- Affects most images
- Degrades the image
- Can cause problems with processing
- Causes?
- Measuring noise: $S/N \text{ ratio} = \frac{\sum_{(i,j)} f^2(i,j)}{\sum_{(i,j)} v^2(i,j)}$
- Correcting noise...
- Types
 - Gaussian
 - Salt and Pepper

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Images Slide 21

21

Colour Images – Other colour spaces

- HSV
- YCrCb
- CIE XYZ
- CIE L*u*v*
- CIE L*a*b*
- Bayer

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Images Slide 20

20

Noise – Salt and Pepper Noise

- Impulse noise
- Noise is maximum or minimum values

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Images Slide 22

22

Noise – Gaussian Noise

- Good approximation to real noise
- Distribution is Gaussian (mean & s.d.)

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Images Slide 23

23

Smoothing – Averaging Filters (linear)

- Linear transformation (convolution)
- Local neighbourhood
 - $f(i,j) = \sum_{(m,n) \in O} h(i-m, j-n) \cdot g(m,n)$
 - Different masks...
 - Local Average
 - Gaussian
- Acceptable results?
 - Suppression of (small) image noise
 - Blurring of edges

$$h = \frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

$$h = \frac{1}{10} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 2 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

$$h = \frac{1}{16} \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$$

```
blur(image,smoothed_image,Size(3,3));
GaussianBlur(image,smoothed_image,Size(5,5),1.5);
```

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Images Slide 25

25

Smoothing

- Removing or reducing noise...
- Linear & non-linear transformations

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Images Slide 24

24

Smoothing examples

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Images Slide 26

26

Smoothing – Median Filter (non-linear)

- Use the median value... 11 18 20 21 23 25 25 30 250
- Not affected by noise Median = 23 Average = 47
- Doesn't blur edges much
- Can be applied iteratively
- Damages thin lines and sharp corners
 - Change region shape
- Computational expensive
 - Standard – $O(r^2 \log r)$
 - Huang – $O(r)$
 - Perreault (2007) – $O(1)$

```
medianBlur(image, smoothed_image, 5);
```

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Images Slide 27

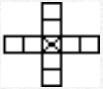


Image Pyramids

- To process images
 - At multiple scales
 - Efficiently
- Technique
 - Smooth image (often Gaussian)
 - Subsample (usually by a factor of 2)

```
pyrDown( image, smaller_image,  
Size( (image1.cols+1)/2, (image1.rows+1)/2 ));
```

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Images Slide 29



Smoothing – Effects of mask size

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Images Slide 28

Learning & Evaluation

- Data / Experience
 - Unsupervised
 - Supervised
- Performance Evaluation
 - Ground Truth
 - Metrics
 - Training vs. Validation

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Learning & Evaluation Slide 1

1

Unsupervised machine learning

- k-means.
 - Data – Unlabelled image data
 - Task – Grouping pixels to reduce the number of colours
 - Performance – Very application/image dependent
 - Measure accuracy or error rate
- Gaussian Mixture Model
 - Data – Unlabelled video data.
 - Task – Identification of moving objects.
 - Performance
 - Accuracy based on hand labelled group truth.

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Learning & Evaluation Slide 3

3

Machine learning

- An algorithm that can learn from data.
 - Data / Experience
 - Some task
 - Often classification
 - Performance improved by data
 - Measure accuracy or error rate
- Data / Experience
 - Unsupervised / Unlabelled.
 - Supervised / Labelled.

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Learning & Evaluation Slide 2

2

Supervised machine learning

- Back-projection.
 - Data – Labelled sample pixel values
 - Task – Classifying pixels as a particular class.
 - Performance
 - Accuracy based on hand labelled group truth.
- Statistical Pattern Recognition / Support Vector Machines
 - Data – Labelled samples (feature vectors).
 - Task – Recognition of objects/shapes.
 - Performance
 - Accuracy based on hand labelled group truth.

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Learning & Evaluation Slide 4

4

1

2

Evaluating Computer Vision Performance

- ✓ Does it work in real time?
 - What is real time?
 - Does it matter?
- ✓ What is the correct answer?
 - Ground truth.
 - How do we generate ground truth?
- ✓ How do we assess how well we have done?
 - Success/failure metrics?
 - Overlap?
 - What about time?

Learning & Evaluation

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 5

5

Performance – Ground Truth

- ✓ What should any computer vision technique achieve?
 - ✓ The correct answer is easy for some techniques. e.g.
 - ✓ Find skin pixels
 - ✓ For some techniques it is harder.
 - ✓ Find the edges.
- ✓ The usefulness of a technique
 - ✓ depends on the application...



Learning & Evaluation

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 7

7

Performance – Ground Truth

- ✓ What should any computer vision system achieve?
 - ✓ The correct answer is easy for some tasks. e.g.
 - ✓ Is this picture me?
 - ✓ Determine the state of the traffic lights.
 - ✓ For some tasks it is harder.
 - ✓ Classify the object in this image.
 - ✓ Count the people.
 - ✓ Drive from A to B.



Learning & Evaluation

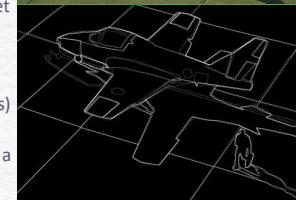
Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 6

6

Ground Truth – Example

- ✓ How should an image be segmented?
 - ✓ Significant edges/regions?
 - ✓ Impossible to get agreement?
- ✓ Berkeley segmentation dataset
 - ✓ 1,000 images
 - ✓ 12,000 segmentations (colour & grayscale)
 - ✓ Hand-labelled (30 subjects)
- ✓ Is this a good way to evaluate a segmentation technique?



From the public "Berkeley Segmentation dataset"

Learning & Evaluation

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 8

8

3

4

Ground Truth – Bounding boxes

- Boxes incorporate parts of other objects/background.
- Labour intensive for video.

CAVIAR:

- Shopping centre (Lisbon)
- Bounding boxes for all people.
- Some have head, hands, and feet annotated.
- Manually annotated.



From the public dataset CAVIAR

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Learning & Evaluation Slide 9

9

Performance – Simple Metrics

- Compute
 - TP
 - TN
 - FP
 - FN
- and then compute...

$$\text{Recall} = \frac{\text{TP}}{\text{TP} + \text{FN}}$$

$$\text{Precision} = \frac{\text{TP}}{\text{TP} + \text{FP}}$$

$$\text{Accuracy} = \frac{\text{TP} + \text{TN}}{\text{Total Samples}}$$

$$\text{Specificity} = \frac{\text{TN}}{\text{FP} + \text{TN}}$$

$$F_\beta = (1 + \beta^2) \cdot \frac{\text{Precision} \cdot \text{Recall}}{(\beta^2 \cdot \text{Precision}) + \text{Recall}}$$

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Learning & Evaluation Slide 11

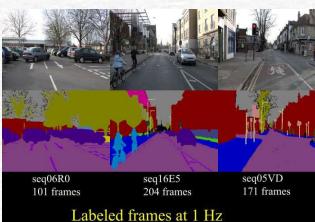
11

Ground Truth – Pixel based

- Better validity.
- Very expensive to create.

CAMVID:

- 32 classes
- 10 minutes @30/15Hz
- 700 frames
- Manually annotated.
- 2nd person inspection.



Labeled frames at 1 Hz
From the public dataset CAMVID

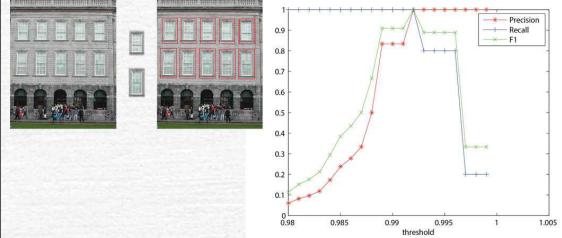
Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Learning & Evaluation Slide 10

10

Performance – Tuning

- We can tune performance by altering parameters
- For example altering the threshold used for accepting a match in a template matching example



Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Learning & Evaluation Slide 12

12

Performance – Precision-Recall Curves

- Alternative (common) visualisation
- Select the parameters which result in the point on the PR curve which is closest to P=1.0, R=1.0

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Learning & Evaluation Slide 13

13

Performance – Confusion Matrix

- It is helpful to know how things are misclassified
- Actual class
- Determined class

		Ground Truth			
		Walking	Running	Standing	Dancing
Result	Walking	0.9	0.0	0.0	0.0
	Running	0.0	0.7	0.1	0.1
Standing	0.1	0.0	0.0	0.0	
Dancing	0.0	0.2	0.0	0.4	
Hopping	0.0	0.1	0.0	0.5	

	A	Not A
A	TP	FP
Not A	FN	TN
Face	3	0
Not Face	1	65532

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Learning & Evaluation Slide 15

15

Performance – Overlap & mAP

- Intersection over Union (IoU)
$$IoU(Detected, GroundTruth) = \frac{\text{Detected} \cap \text{GroundTruth}}{\text{Detected} \cup \text{GroundTruth}}$$

- Values of IoU > 0.5 are often considered good matches.
- When considering multiple classes...
- Mean Average Precision (mAP)
$$mAP = \frac{1}{\text{classes}} \sum_{c \in \text{classes}} \frac{\#TP(c)}{\#TP(c) + \#FP(c)}$$

ResNet result on COCO dataset (2015)

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Learning & Evaluation Slide 14

14

Training & Validation

- Typically datasets are divided into two parts:
- Training set.
- Validation/Testing set.
- This can be
- Static
- Random
- Leave p-out cross validation
- Leave p samples out – for testing. Train on the rest.
- Do multiple rounds of training and testing using different partitions.
- Average results.
- Intended to determine how the results will generalise.

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Learning & Evaluation Slide 16

16

Outcome

- So what does testing tell me?
 - How well the system/technique works on the test data (images/videos)...
 - Actually, how close the system/technique gets to the manually created ground truth...
- How representative is the test set?
 - What conditions matter?
 - Can I ever be sure that a system will work in general?

Learning & Evaluation

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 17

Recognition

- Template Matching
 - Chamfer Matching
- Statistical Pattern Recognition (SPR)
- Support Vector Machines (SVM)
- Cascade of Haar classifiers (Haar)
- Principal Components Analysis (PCA)

Recognition

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 1

Summary

- Learning (Data – Task – Performance)
 - Unsupervised
 - Supervised
- Ground Truth
 - Expensive to obtain.
 - Often unclear what is the right answer.
- Metrics
 - Precision, Recall, Accuracy
 - Tuning
 - IoU, mAP
- Training vs. Validation
 - Leave p-out cross validation

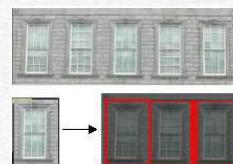
Learning & Evaluation

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 18

Template Matching - Topics

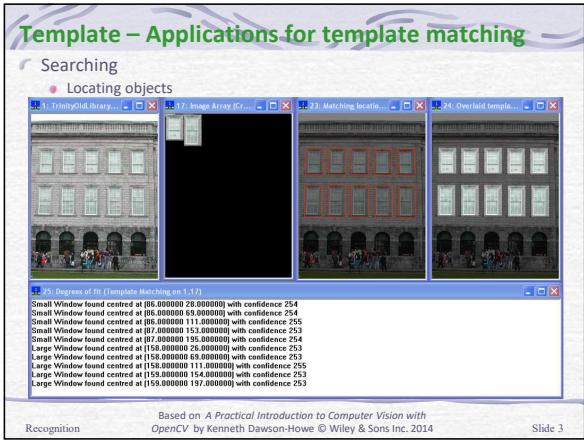
- Applications
- Matching criteria
- Control strategies
- Chamfer Matching



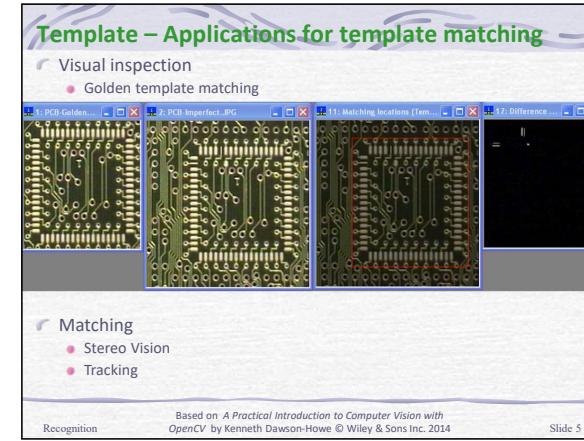
Recognition

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

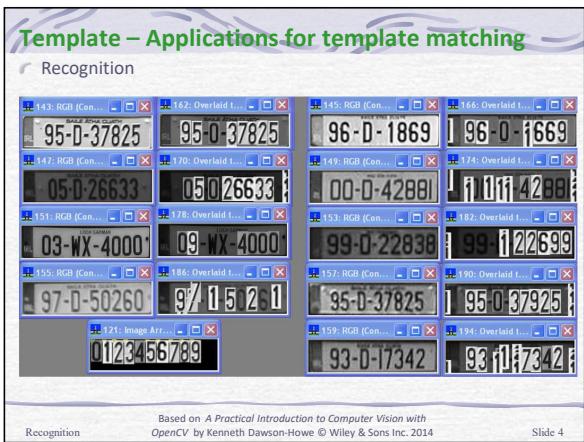
Slide 2



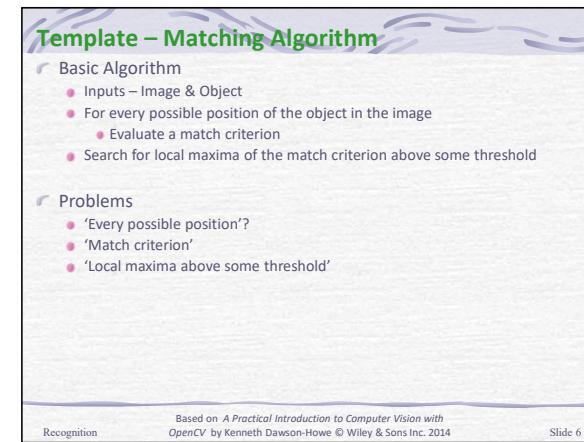
3



5



4



6

Template – Matching criteria

$$D_{SquareDifferences}(i,j) = \sum_{(m,n)} (f(i+m, j+n) - t(m, n))^2$$

$$D_{NormalisedSquareDifferences}(i,j) = \frac{\sum_{(m,n)} (f(i+m, j+n) - t(m, n))^2}{\sqrt{\sum_{(m,n)} f(i+m, j+n)^2 \sum_{(m,n)} t(m, n)^2}}$$

$$D_{CrossCorrelation}(i,j) = \sum_{(m,n)} f(i+m, j+n) \cdot t(m, n)$$

$$D_{NormalisedCrossCorrelation}(i,j) = \frac{\sum_{(m,n)} f(i+m, j+n) \cdot t(m, n)}{\sqrt{\sum_{(m,n)} f(i+m, j+n)^2 \sum_{(m,n)} t(m, n)^2}}$$

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Recognition Slide 7

7

Template Matching in OpenCV

```
Mat matching_space;
matching_space.create(
    search_image.cols-template_image.cols+1,
    search_image.rows-template_image.rows+1, CV_32FC1 );
matchTemplate( search_image, template_image,
    matching_space, CV_TM_CCORR_NORMED );
// Other measures: CV_TM_CCORR, CV_TM_SQDIFF,
// CV_TM_SQDIFF_NORMED
```

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Recognition Slide 9

9

Template – Matching criteria example

- Simple example
- Matching criteria
- Boundaries...

0	0	0	0	0	1	0	0
0	0	0	0	0	1	0	0
0	0	0	1	1	1	1	0
0	0	0	1	1	0	1	0
0	0	0	1	0	0	1	0
0	0	0	1	0	0	1	0
0	0	0	1	0	1	1	0
0	0	0	0	0	0	0	0

Image

1	1	1	1	1
1	0	0	1	1
1	0	0	1	1
1	0	0	1	1
1	1	1	1	1

Template

11	13	14	11	14
10	14	16	8	14
9	12	12	1	10
10	14	15	7	15

Matching Space

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Recognition Slide 8

8

Template – Finding Local Maxima in OpenCV

- Local maxima – Dilate + Look for unchanged values + Threshold
- Local minima – Erode + Look for unchanged values + Threshold

```
Mat dilated, thresholded_matching_space, local_maxima,
thresholded_8bit;
dilate( matching_space, dilated, Mat());
compare( matching_space, dilated, local_maxima, CMP_EQ );
threshold( matching_space, thresholded_matching_space,
           threshold, 255, THRESH_BINARY );
thresholded_matching_space.convertTo( thresholded_8bit,
                                     CV_8U );
bitwise_and( local_maxima, thresholded_8bit, local_maxima );
```

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Recognition Slide 10

10

Template – Control Strategies for Matching

- Goal: Localise close copies
 - Size, orientation
 - Geometric distortion
- Use an image hierarchy
 - Low resolution first
 - Limit higher resolution search
- Search higher probability locations first
 - Known / learnt likelihood
 - From lower resolution

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Recognition Slide 11

11

Template – Chamfering

- Compute chamfered image for a binary edge image.

```

For every point
  If (edge point)
    Set c(i, j) = 0
    Else Set c(i, j) = ∞
  For j = min to max
    For i = min to max
      c(i, j) = minq∈AL [distance( (i, j), q ) + f( q )]
  For j = max to min
    For i = max to min
      c(i, j) = minq∈BR [distance( (i, j), q ) + f( q )]

```

Object pixels (zeros)

∞	∞	∞	∞	∞	∞	∞
∞	∞	∞	∞	∞	∞	∞
∞	0	0	0	0	0	0
∞	0	0	0	0	0	0
∞	0	0	0	0	0	0
∞	0	0	0	0	0	0
∞	0	0	0	0	0	0
∞	0	0	0	0	0	0
∞	0	0	0	0	0	0

After the first stage of processing

3.8	2.8	2.4	2	2	1	0	1	2
3.4	2.4	1.4	1	1	0	1	1.4	
3	2	1	0	0	0	0	1	
3	2	1	0	0	1	0	1	
3	2	1	0	1	1	0	1	
3	2	1	0	1	1.4	0	1	
3	2	1	0	0	0	0	1	
3	2	1	1	1	1	1	1	

Canny(gray_image, edge_image, 100, 200, 3);
threshold(edge_image,edge_image,127,255,THRESH_BINARY_INV);
distanceTransform(edge_image,chamfer_image,CV_DIST_L2,3);

Chamfer Image

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Recognition Slide 13

13

Template – Chamfer matching

- Template matching requires very close matches
- Objects often appear very slightly different
 - Orientation
 - Noise
 - Sampling
- We want a more flexible approach

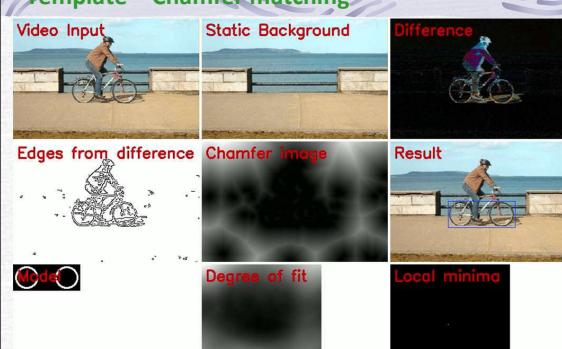


Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Recognition Slide 12

12

Template – Chamfer matching



Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Recognition Slide 14

Template – Chamfering for matching

- Compare with boundary template
 - Sum of chamfer values for boundary points
 - Low values best

Chamfer matching is a simple routine to write (provided in the text);

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

15

SPR – Features: Minimum Bounding Rectangle

- Turn rectangle through discrete steps
- Only one quadrant
- Metrics:
 - Length to Width ratio
 - Length / Width
 - Rectangularity
 - Area / (Length * Width)
 - Convex Hull to Minimum Bounding Rectangle area ratio
 - Area inside convex hull / (Length * Width)

`RotatedRect min_bounding_rectangle = minAreaRect(contours[contour_number]);`

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

17

SPR – Statistical Pattern Recognition

- Probability Review
- Features
- Classification

	Rectangularity	Elongatedness
Albatross	0.28	3.63
Vulture	0.45	5.56
Hawk	0.39	2.39
Falcon	0.33	2.11

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

16

SPR – Features: Elongatedness

- CANNOT BE the ratio of the length and width of the minimum bounding rectangle
- Ratio of region area divided by the square of it's thickness

$$\text{Elongatedness} = \frac{\text{area}}{(2d)^2}$$

	Rectangularity	Elongatedness
Circle	0.79	1.00
Square	1.00	1.00
Rectangle	1.00	1.85
Circle	1.00	1.82
Square	0.99	1.82
Rectangle	0.39	1.82
Circle	0.40	1.00
Square	1.00	1.00

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

18

SPR – Probability Review – Basics

- $P(A) = \lim_{n \rightarrow \infty} N(A)/n$
- Probability of two events A and B:
 - Independent: $P(AB) = P(A)P(B)$
 - Dependent: $P(AB) = P(A|B)P(B)$
 - Conditional Probability $P(A|B)$
- Typical problem:
 - Given some evidence x from an unknown object.
 - What class W_i is the object?
 - Training
 - A-priori probability – $p(x | W_i)$
 - Relative probability – $p(W_i)$
 - A-posteriori probability – $p(W_i | x)$

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Recognition Slide 19

19

SPR – Probability Density Functions

- Given a class (W_i) and an event (x)
 - Determine the probability of any particular value occurring...

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Recognition Slide 21

21

SPR – Probability Review – Bayes Theorem

- For two classes A and B the a-posteriori probability is:

$$P(B|A) = P(A|B)P(B) / P(A)$$
- Where W_i forms a partitioning of the event space:

$$p(W_i | x) = \frac{p(x | W_i)P(W_i)}{\sum_j p(x | W_j)P(W_j)}$$

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Recognition Slide 20

20

10

SPR – Classification

- Object recognition
 - Classes (w_1, w_2, \dots, w_R)
- Classifier
 - Input Pattern / features (x_1, x_2, \dots, x_n)
- Feature space
 - Choosing the features ([Example](#))
 - Clusters in feature space
- Separability
 - Hyper-surfaces
 - Linear separability
 - Inseparable classes
- Probabilistic Classifier

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Recognition Slide 22

22

11

SPR – Probabilistic classifier

- Optimal classifier
- Based on Probability Density Functions
 - Remove normalisation from Bayes rule... $\sum_j p(x|W_j)P(W_j)$
 - Resultant Discrimination function: $g_r(x) = p(x|W_r)P(W_r)$
 - $g_r(x) \geq g_s(x)$
 - Unknown class: $g_r(x) > \text{threshold}$
- Mean Loss Function: $J(q) = \int_x \sum_{s=1..R} \lambda[d(x, q)|W_s]p(x|W_s)P(W_s) dx$
where $\lambda[W_r|W_s] = \begin{cases} 0 & \text{if } r = s \\ 1 & \text{otherwise} \end{cases}$
- Advantages / Disadvantages
 - Accuracy
 - Extensive training

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Recognition Slide 23

23

SPR – Features: Convex hull

```
vector<vector<Point>> hulls(contours.size());
for (int contour=0; (contour<contours.size()); contour++)
{
    convexHull(contours[contour], hulls[contour]);
}
```

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Recognition Slide 25

25

SPR – Features: Convex hull

- Smallest convex region
- Algorithm
 - Start with
 - Any point on convex hull
 - Previous vector direction
 - Search all other boundary points
 - Find point with least angle to previous vector
 - Switch to new point and vector
 - Go to 2 unless new point = start point.

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Recognition Slide 24

24

SPR – Features: Concavities and Holes

- Identify concavities from the convex hull
- Identify holes in the binary shape

Shape	Hull	Height	Width	# holes	Hole Area	Concavity Details	
9	2.62	0.81	1	1.8	3.29@137	0.03@66	
5	2.54	0.85	0	4	3.0@140	2.4@302	
0	2.54	0.87	1	4.5	3.0@171	0.1@116	
3	2.62	0.81	0	5	7.6@144	0.4@343	
7	2.62	0.81	0	1	2.0@137	0.0@100	
8	2.62	0.81	2	1.6	1.5	3.0@354	0.3@203
2	2.43	0.84	0	2	3.9@201	2.8@40	
5	2.62	0.83	0	4	3.0@143	2.4@300	

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Recognition Slide 26

26

SPR – Features: Concavities and Holes

```
vector<vector<int>> hull_indices(contours.size());
vector<vector<Vec4i>> convexity_defects(contours.size());
for (int contour=0; (contour<contours.size()); contour++)
{
    convexHull( contours[contour], hull_indices[contour] );
    convexityDefects( contours[contour], hull_indices[contour],
                      convexity_defects[contour]);
}
```

Recognition

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 27

SPR – Classifier learning

Training set

- Must be representative
- Must be inductive
- Can use Kernel Density Estimation

Training set size

- Training set provides the unknown statistical information
- Size will typically have to be increased several times

Learning strategies

- Supervised:
 - Probability density estimation – estimating $p(x|w_r)$ & $P(w_r)$
 - Training set includes class specification for every instance
- Unsupervised:
 - Cluster Analysis
 - Look for similarities in feature space

Recognition

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 29

SPR – Features: Perimeter Length and Circularity

- Perimeter length
- Approximately the number of boundary elements
`contours[contour].size()`
- Should really take into account direction
- Circularity = Perimeter length / ($4 * \pi * \text{Area}$)



Shape	Perimeter	Area	Bounding Box	Convex Hull	Hu Moments	Hole Areas
0	80	651	553	553	1.7	0.00
1	460	1	460	460	0.00	0.00
2	111	272	272	272	0.44	0.00
3	111	272	272	272	0.44	0.00
4	121	351	522	530	0.40	0.06
5	347	592	442	420	0.01	0.00
6	122	351	570	504	0.43	0.06
7	102	399	620	534	0.23	0.01
8	94	267	551	394	0.46	0.09
9	63	438	589	589	0.23	0.01
	9	101	389	589	0.23	0.01
					0.00	74

Recognition

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 28

SPR: Real example



Shape	Height / Width	# holes	Hole Areas	Concavities	Concavity Details
0	2.46	0.81	1	4.1	0
1	4.71	0.58	0	0	0.4@243 0.03@261
2	2.83	0.81	0	5	3.6@213 2.5@45
3	2.83	0.81	0	3	5.7@185 0.2@354
4	2.38	0.61	1	0.6	3.0@9101 0.1@207
5	2.54	0.81	0	4	3.1@140 2.4@302
6	2.13	0.79	1	1.5	3.1@317 0.03@94
7	2.62	0.81	1	0.4	3.4@333 0.1@241
8	2.83	0.81	2	1.81	2.0@205 0.2@355
9	2.62	0.81	1	1.9	4.0@143 0.1@240

Recognition

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 30

Support Vector Machines

- Popular technique for two-class problem
- Consider two linearly separable classes in n -dimensional feature space
- Many hyper-surfaces usually exist
- Optimal classification
 - Maximise the margin between the classes.

Recognition © Kenneth Dawson-Howe 2015 Slide 31

31

Support Vector Machines – Formulation

- Given many n -dimensional training samples x_i with associated class identifiers $\omega_i = \{-1, 1\}$
- Define the separating hyper-planes $w \cdot x + b = 0$
 $w \cdot x + b = 1 \quad w \cdot x + b = -1$
- Constraint: $\omega_i(w \cdot x_i + b) \geq 1$
- If x^+ is a point on $w \cdot x + b = 1$ and x^- is the nearest point on $w \cdot x + b = -1$
 - then $\lambda w = x^+ - x^-$
- Multiply by w $\lambda w \cdot w = x^+ \cdot w - x^- \cdot w = (1 - b) - (-1 - b) = 2$
 - Therefore $\lambda w \cdot w = \lambda \|w\|^2 = 2$ as $\|w\| = \sqrt{w \cdot w}$ and...
- $\|x^+ - x^-\| = \|\lambda w\| = \left\| \frac{2 \cdot w}{\|w\|^2} \right\| = \frac{\|2 \cdot w\|}{\|w\|} = \frac{\|2 \cdot w \cdot w\|}{\|w\|} = \frac{2}{\|w\|}$
- Hence we must minimise $\|w\|$ to maximise the separation, subject to $\omega_i(w \cdot x_i + b) \geq 1$

Recognition © Kenneth Dawson-Howe 2015 Slide 33

33

Support Vector Machines – Example

From Vision-Based Traffic Data Collection Sensor for Automotive Applications by Llorente DF, Sánchez S, Ocaña M, Sotelo MA - (2010) http://openi.nlm.nih.gov/detailedresult.php?img=3270873_sensors-10-00860f12&rpp=4

Recognition Slide 32

32

Support Vector Machines – Lagrangian Optimisation

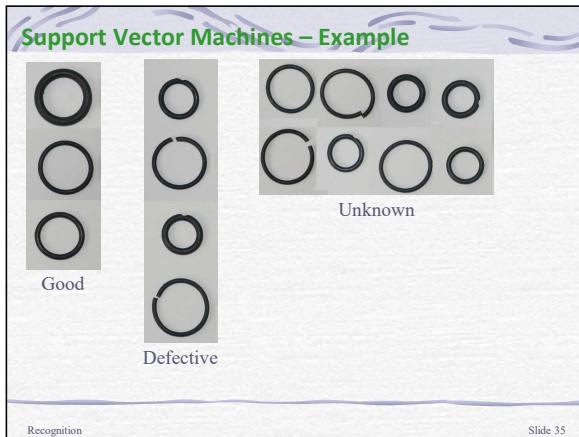
- This is a minimisation problem subject to constraints which is amenable to solution by Lagrangian optimisation

$$L(w, b, \alpha) = \frac{\|w\|^2}{2} - \sum_{i=1}^n \alpha_i [\omega_i (w \cdot x_i + b) - 1]$$

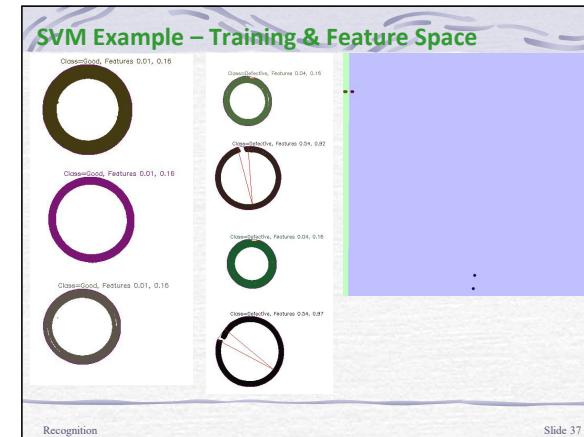
- To find the minimum we set the partial differentials to 0
- Classification is simply based on
 - entering an unknown feature vector x_i
 - into the original equation $f(x_i) = w \cdot x_i + b$
 - and classifying based on whether the value is positive or negative (preferably ≥ 1).

Recognition © Kenneth Dawson-Howe 2015 Slide 34

34



35



37

SVM Example – Training

```
Ptr<SVM> svm = ml::SVM::create();
svm->setType(SVM::C_SVC);
svm->setKernel(SVM::LINEAR);
Mat labelsMat(number_of_samples, 1, CV_32SC1, labels);
Mat trainingDataMat(number_of_samples, 2, CV_32FC1,
                     training_data);
Ptr<ml::TrainData> tData = ml::TrainData::create(trainingDataMat,
                                                 ml::SampleTypes::ROW_SAMPLE, labelsMat);
svm->train(tData);
```

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Recognition

Slide 36

36

SVM – Support vectors

```
Mat support_vectors = svm->getSupportVectors();
for (int support_vector_index = 0; support_vector_index <
support_vectors.rows; ++support_vector_index)
{
    const float* v = support_vectors.ptr<float>(support_vector_index);
    circle(feature_space, Point((int)v[0],(int)v[1]), 3, Scalar(0,0,255));
}
```

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Recognition

Slide 38

38

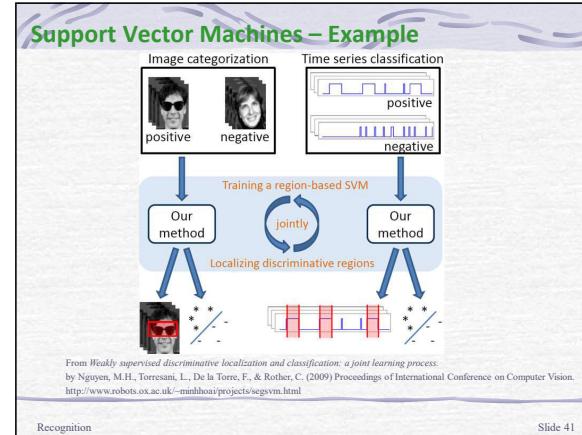
SVM Example – Prediction

```
// Try to predict the class
Mat sampleMat = (Mat<float>(1,2) << feature[0], feature[1]);
float prediction = svm->predict(sampleMat);
class_id = (prediction > 0.0) ? 1 : (prediction < 0.0) ? 2 : 0;
```

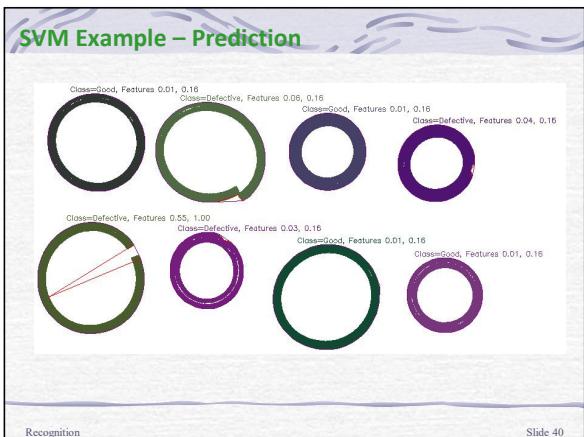
Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Recognition Slide 39

39

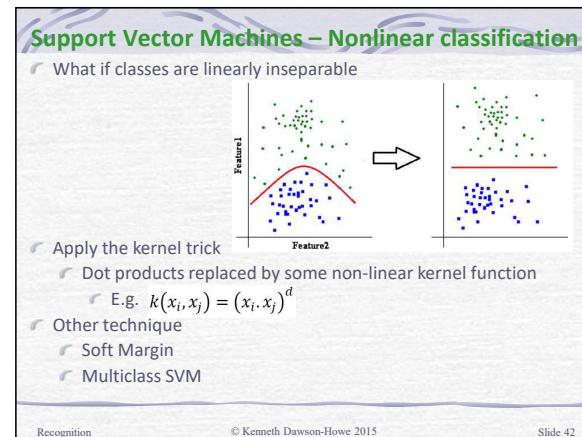


41



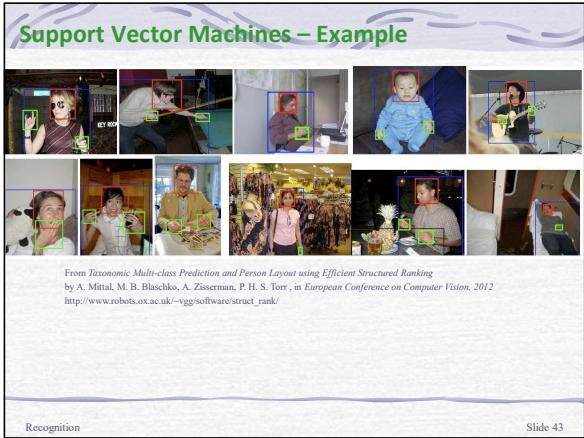
40

20

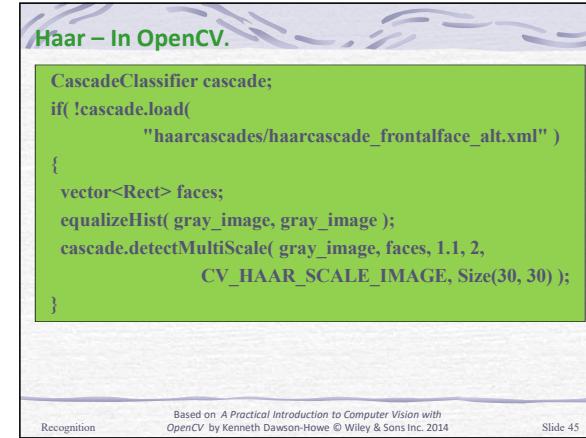


42

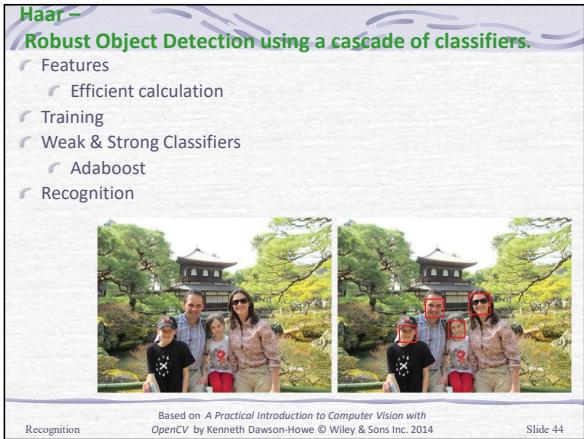
21



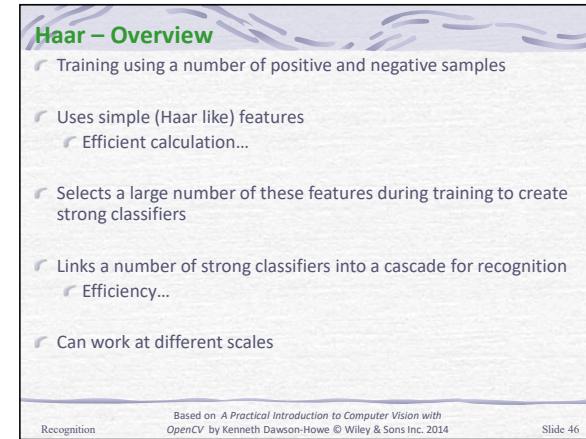
43



45



44



46

Haar – Features

Features determined as the difference of the sums of a number of rectangular regions

- Place the mask in a specific location and at a specific scale
- Then subtract the sum of the 'white pixels' from the sum of the 'black pixels'
- Why does this work?

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Recognition Slide 47

47

Haar – Training

- Number of possible features.
 - the variety of feature types
 - allowed variations in size and position
- Training must
 - Identify the best features to use at each stage
 - To do this positive and negative samples are needed...

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Recognition Slide 49

49

Haar – Efficient Calculation – Integral Image

- Integral Image:
 - Every point $ii(i,j) = \sum_{i'=0,i} \sum_{j'=0,j} \text{image}(i',j')$
 - Sum of points in rectangle D:

$$\text{sum}(D) = ii(p) + ii(s) - ii(q) - ii(r)$$

- Features can be computed at any scale for the same cost

An **integral** function is provided which computes the integral image both normally and at 45°

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Recognition Slide 48

48

Haar – Weak & Strong Classifiers

- Weak Classifier
 - $p_j \text{feature}_j(x) < p_j \vartheta_j$
 - Tune threshold (ϑ_j)
- Strong Classifiers
 - Combine a number of weak classifiers...
 - E.g. 100% true positives with only 40% false positives using only two face features...

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Recognition Slide 50

50

Haar – Strong Classifiers – AdaBoost

Given n example images x_1, x_n together with classifications y_1, y_n
where $y_i = 0, 1$ for negative and positive examples respectively.

Initialise weights $w_{t,i} = 1 / (2 * (m * (1 - y_j) + l * y_j))$
where m and l are the number of negative and positive examples respectively.

For $t=1, \dots, T$

1. Normalize the weights (i.e. for all i): $w_{t,i} = w_{t,i} / (\sum_{j=1..n} w_{t,j})$
2. For each feature, j , train a weak classifier $h_j(x)$ and evaluate the error taking into account the weights: $\epsilon_j = \sum_i w_{t,i} |h_j(x_i) - y_i|$
3. Select the classifier, $h_j(x)$, with the lowest ϵ_j , save as $c_t(x)$ with error E_t
4. Update the weights (i.e. for all i): $w_{t+1,i} = w_{t,i} \beta_t^{(1-e)}$
where $e_t = |c_t(x_i) - y_i|$ and $\beta_t = E_t / (1-E_t)$

The final strong classifier is: $h(x) = 1 \text{ if } \sum_{t=1..T} \alpha_t c_t(x) \geq \frac{1}{2} \sum_{t=1..T} \alpha_t$,
0 otherwise
where $\alpha_t = \log 1/\beta_t$

Recognition Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014 Slide 51

51

Haar – Recognition

- Face recognition
- 38 stages
- 6000+ features
- 4916 positive samples
- 9544 negative samples
- Scale independence

Recognition Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014 Slide 53

53

Haar – Classifier cascade

- Object recognition is possible with a single strong classifier
- To improve detection rates AND to reduce computation time:
 - A cascade of strong classifiers can be used
 - Each stage either accepts or rejects
 - Only those accepted pass to the next stage
 - Efficient computation...
- Strong classifiers trained using AdaBoost
 - On the remaining set of data

Recognition Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014 Slide 52

52

Principal Component Analysis (PCA)

- Statistical Technique for
 - Data analysis
 - Data compression
 - Recognition
 - Analyses data covariance
 - Identifies principal directions
- For example:
 - Given 2 dimensional data
 - and N samples/vectors
 - Find
 - the mean sample
 - the direction of maximum covariance
 - the direction orthogonal to this

Recognition © Kenneth Dawson-Howe, 2015 Slide 54

54

PCA – Background – Eigenvalues and Eigenvectors

- Consider a square matrix A
- The eigenvalues of A are the roots of the *characteristic equation*: $\text{determinant}(A - \lambda I) = |A - \lambda I| = 0$
- For each eigenvalue λ there will be an eigenvector x such that $Ax = \lambda x$
- Many possible values of x
- An $n \times n$ matrix will have n eigenvalues
- Consider $n=2$. There will be 2 eigenvalues: λ_1, λ_2
 - with corresponding eigenvectors x_1, x_2
 - where $Ax_1 = \lambda_1 x_1$ and $Ax_2 = \lambda_2 x_2$

Recognition

© Kenneth Dawson-Howe, 2015

Slide 55

55

PCA – Theory – Organise Data

- In PCA we have N samples/vectors in some n -dimensional space
 - Combine into a data matrix D
 - Each row is a sample
 - Determine the mean of the samples $D = \begin{bmatrix} x_1^1 & x_2^1 & \dots & x_n^1 \\ x_1^2 & x_2^2 & \dots & x_n^2 \\ \vdots & \vdots & & \vdots \\ x_1^N & x_2^N & \dots & x_n^N \end{bmatrix}$
 - Compute mean-centred data U
- $$\mu = \begin{bmatrix} \sum_{i=1}^N x_1^i / N & \sum_{i=1}^N x_2^i / N & \dots & \sum_{i=1}^N x_n^i / N \end{bmatrix}$$
- $$U = D - \begin{bmatrix} \mu \\ \mu \\ \vdots \\ \mu \end{bmatrix} = \begin{bmatrix} x_1^1 - \sum_{i=1}^N x_1^i / N & x_2^1 - \sum_{i=1}^N x_2^i / N & \dots & x_n^1 - \sum_{i=1}^N x_n^i / N \\ x_1^2 - \sum_{i=1}^N x_1^i / N & x_2^2 - \sum_{i=1}^N x_2^i / N & \dots & x_n^2 - \sum_{i=1}^N x_n^i / N \\ \vdots & \vdots & & \vdots \\ x_1^N - \sum_{i=1}^N x_1^i / N & x_2^N - \sum_{i=1}^N x_2^i / N & \dots & x_n^N - \sum_{i=1}^N x_n^i / N \end{bmatrix}$$

Recognition

© Kenneth Dawson-Howe, 2015

Slide 57

57

PCA – Background – Eigenvalues and Eigenvectors

- Combining we get $A[x_1 \ x_2] = [x_1 \ x_2] \begin{bmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{bmatrix}$
 - and $A\Phi = \Phi\Lambda$
 - where $\Phi = [x_1 \ x_2]$ and $\Lambda = \begin{bmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{bmatrix}$
- If we normalise the eigenvectors then $\Phi\Phi^T = \Phi^T\Phi = I$
 - so $\Phi^T A \Phi = \Phi^T \Phi \Lambda = \Lambda$
 - and $A = A\Phi\Phi^T = \Phi\Lambda\Phi^T$

Recognition

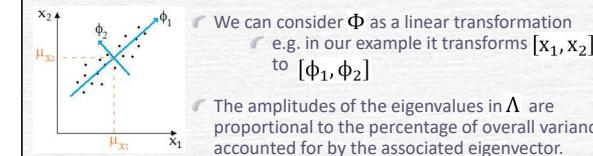
© Kenneth Dawson-Howe, 2015

Slide 56

56

PCA – Theory – Covariance matrix

- Using the mean-centred data U , compute the covariance matrix $\Sigma = U^T U / (N - 1)$
- We can then determine $\Phi^T \Sigma \Phi = \Lambda$
 - where the eigenvectors are in an orthogonal matrix Φ
 - and the eigenvalues are in an ordered diagonal matrix Λ

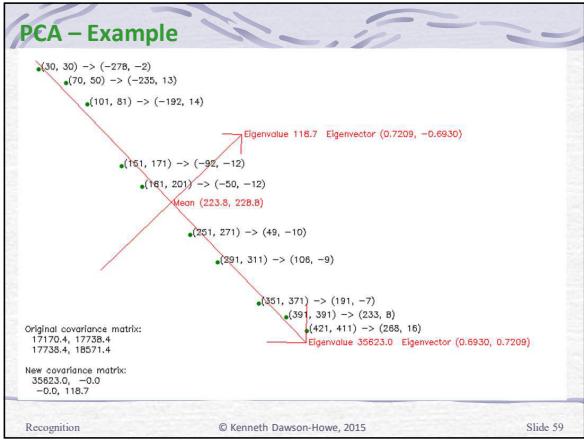


Recognition

© Kenneth Dawson-Howe, 2015

Slide 58

58



PCA – Basics

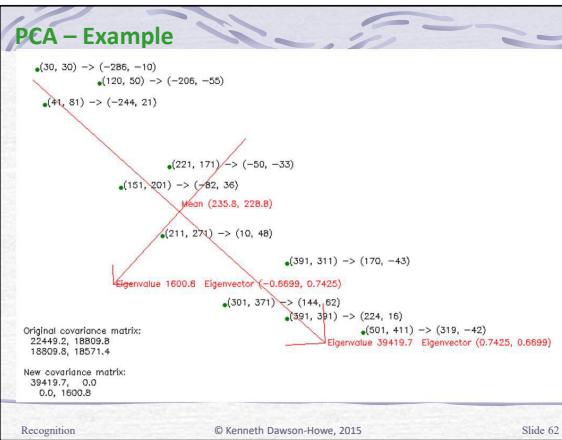
```
Mat transformed_samples_matrix = samples_matrix.clone();
for (int sample_no=0; sample_no<number_of_samples; sample_no++)
{
    Mat sample = samples_matrix.row(sample_no);
    Mat transformed_sample;
    pca.project(sample, transformed_sample);
    transformed_sample.row(0).copyTo(
        transformed_samples_matrix.row(sample_no));
}
Mat covariance,average,new_covariance,new_average;
calcCovarMatrix(samples_matrix,covariance,average,
    CV_COVAR_NORMAL | CV_COVAR_ROWS);
calcCovarMatrix(transformed_samples_matrix,new_covariance,
    new_average,CV_COVAR_NORMAL | CV_COVAR_ROWS);
```

Recognition © Kenneth Dawson-Howe 2015 Slide 61

PCA – Basics

```
float samples[][2] = { {30,30}, {70,50}, ..., {421,411}, {391,391} };
int number_of_samples = sizeof(samples)/sizeof(int[2]);
Mat samples_matrix(number_of_samples, 2, CV_32FC1, samples);
...
PCA pca(samples_matrix, Mat(), 0, 2 );
Mat eigenvalues = pca.eigenvalues;
Mat eigenvectors = pca.eigenvalues;
Mat mean = pca.mean;
```

Recognition © Kenneth Dawson-Howe 2015 Slide 60



PCA – Compression

- If the variance on an axis is very small
 - Variance caused by noise?
 - Drop this dimension?

Recognition © Kenneth Dawson-Howe, 2015 Slide 63

63

PCA – Face Recognition

- Common example application of PCA
- Each face image is treated as a vector
 - with $n=\text{rows} \times \text{columns}$ dimensions!

Recognition © Kenneth Dawson-Howe, 2015 Slide 65

65

PCA – Example

Original covariance matrix:
17617.8, 17617.8
17617.8, 17617.8

New covariance matrix:
35235.5, 0.0
0.0, 0.0

(30, 30) -> (-275, 0)
(50, 50) -> (-247, 0)

(101, 101) -> (-175, 0)
(151, 151) -> (-104, 0)

(201, 201) -> (-34, 0)
Mean (224.8, 224.8)
(251, 251) -> (37, 0)

(311, 311) -> (122, 0)
(351, 351) -> (178, 0)

(391, 391) -> (235, 0)
(411, 411) -> (253, 0)

Eigenvalue 0.0 Eigenvector (0.7071, -0.7071)
Eigenvalue 35235.5 Eigenvector (0.7071, 0.7071)

Recognition © Kenneth Dawson-Howe, 2015 Slide 64

64

32

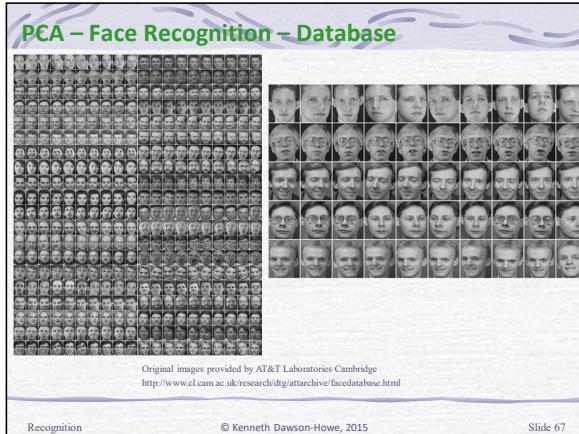
PCA – Face Recognition

- Create a matrix D
 - where each face image is a row
 - D is a $N \times n$ matrix
- Create mean centred data U
 - Determine the mean of the rows (face images)
 - Subtract from all of the rows
- Compute the covariance matrix $\Sigma = U^T U / (N - 1)$
- Solve for Eigenvectors and Eigenvalues as before.
- Normalise the Eigenvectors
- Select m Eigenvectors (with the largest m Eigenvalues)
 - Usually 20-50 for face recognition.
- For an unknown image find its location in PCA space
 - Look for the closest match
$$p_\phi = (p_x - \mu_x) \cdot \Phi_{pca}^T$$

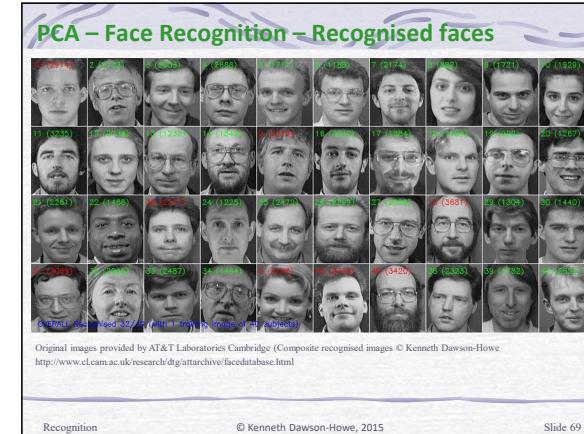
Recognition © Kenneth Dawson-Howe, 2015 Slide 66

66

33



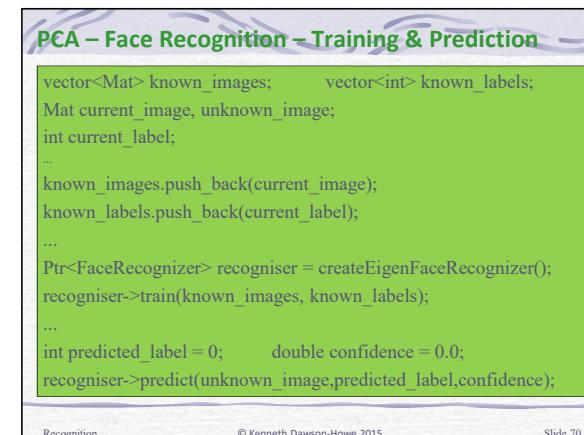
67



69



68



70

PCA – Image Reconstruction using Eigenvectors

- We can reconstruct images from PCA space

$$p_x = p_{\phi} \cdot \Phi_{pca}^T + \mu_x$$
- Using some number (m) of principal components
- So that some percentage of variance is accounted for

$$\% \text{ of Total Variance} = \frac{\sum_{i=1}^m \lambda_i}{\sum_{j=1}^n \lambda_j} \times 100.$$

Figure 3.23: 32 original images of a boy's face, each 321×261 pixels. © Census Learning 2015.

Figure 3.24: Reconstruction of the image from four basic vectors \mathbf{b}_i , $i = 1, \dots, 4$ which can be displayed as images. The linear combination was computed as $q_1 \mathbf{b}_1 + q_2 \mathbf{b}_2 + q_3 \mathbf{b}_3 + q_4 \mathbf{b}_4 = 0.078 \mathbf{b}_1 + 0.062 \mathbf{b}_2 - 0.182 \mathbf{b}_3 + 0.179 \mathbf{b}_4$. © Census Learning 2015. © Cengage Learning 2015. Reproduced with permission from *Image Processing, Analysis and Machine Vision* by Milan Sonka, Vaclav Hlavac and Roger Boyle.

Recognition

Slide 71

71

Region Segmentation

- Binary Regions
- k means clustering
- Watershed Segmentation
- Mean Shift Segmentation

Regions Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014 Slide 1

1

Segmentation

- Split images into smaller parts
 - Preferably corresponding to (parts of) objects.
- Two main approaches
 - Region based
 - Edge based
- Video segmentation
 - Breaking video into clips
 - Segmenting objects/regions consistently over time

Regions Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014 Slide 2

2

Connectivity – Paradoxes

- Use pixel Adjacency to build contiguous regions
 - Objects
 - Background
 - Holes

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Regions Slide 3

3

Connectivity –

- What do we actually want?

- One possibility
 - Treat background using 4-adjacency
 - Treat object using 8-adjacency
 - Treat holes using 4-adjacency
 - Treat objects in holes using 8-adjacency
 - ...

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Regions Slide 5

5

Connectivity – 4 adjacency & 8 adjacency

- Have to use either 4-adjacency or 8-adjacency
 - Label each non-zero pixel...

3	2	1
4	8	0
5	6	7

3	2	1
4	8	0
5	6	7

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Regions Slide 4

4

2

Connectivity – Connected Components Analysis

- Search image row by row
 - Label each non-zero pixel
 - If previous pixels are all background
 - Assign New Label
 - Otherwise
 - Pick any label from the previous pixels
 - If any of the other previous pixels have a different label
 - Note equivalence
- Relabel equivalent labels.

Previous pixels:

■	■	■
■	■	■
■	■	■

8-connectivity 4-connectivity

Equivalences:

■	■
■	■

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Regions Slide 6

6

3

Connectivity – Extracting regions

Regions

```
00000000000000000000
JJJJJJJJJJJJJJJJJJJJ
WWWWWWWWWWWWWWWWWW
BABABABABABABABABAB
```

```
vector<vector<Point>> contours;
vector<Vec4i> hierarchy;
findContours( binary_image, contours, hierarchy,
              CV_RETR_TREE, CV_CHAIN_APPROX_NONE );
```

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 7

7

k means Clustering

We'd like to

- identify significant colours in images
 - Concise descriptions
 - Object tracking
- reduce the number of colours in any image
 - Compression

Regions

How do we find the best colours?

k means clustering

- Creates k clusters of pixels
- Unsupervised learning

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 9

9

Connectivity – Labelling regions

Regions

```
00000000000000000000
JJJJJJJJJJJJJJJJJJJJ
WWWWWWWWWWWWWWWWWW
BABABABABABABABABAB
```

```
for (int contour=0; (contour < contours.size()); contour++)
{
  Scalar colour( rand()&0xFF,rand()&0xFF,rand()&0xFF );
  drawContours( contours_image, contours, contour, colour,
                CV_FILLED, 8, hierarchy );
}
```

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 8

8

k means Clustering – Algorithm

Number of clusters (k) is known in advance (or determine the k with the maximum confidence)

Initialise the k cluster exemplars either randomly or use the first k patterns or ...

1st pass: Allocate patterns to the closest existing cluster exemplar and recompute the exemplar as the centre of gravity

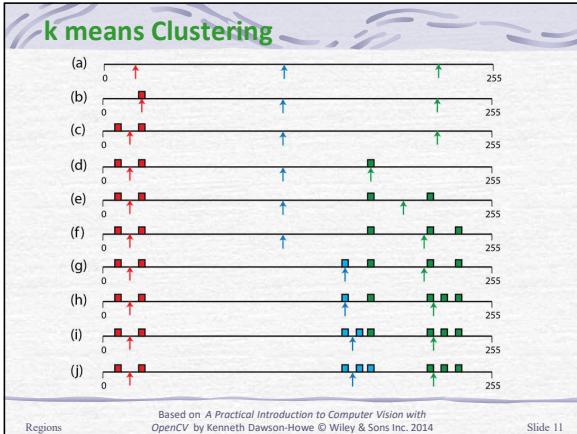
2nd pass: Using the final exemplars from the first pass allocate all patterns cluster exemplars.

Regions

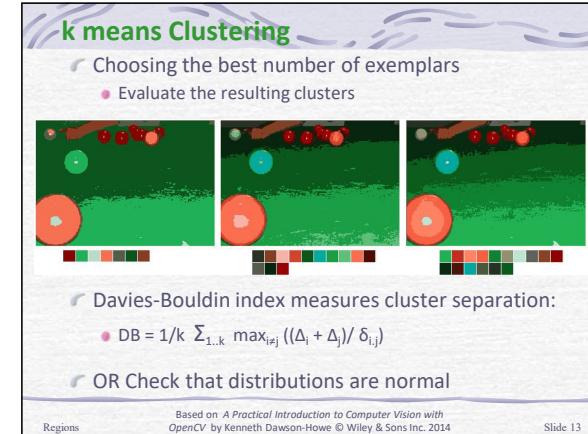
Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 10

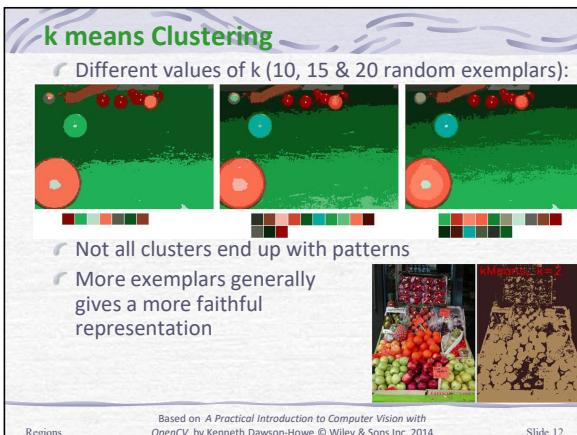
10



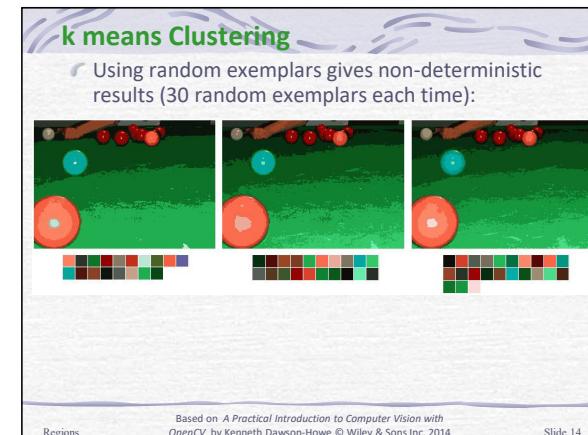
11



13



12



14

k means Clustering

```
// Store the image pixels as an array of samples
Mat samples(image.rows*image.cols, 3, CV_32F);
float* sample = samples.ptr<float>(0);
for(int row=0; row<image.rows; row++)
    for(int col=0; col<image.cols; col++)
        for (int channel=0; channel < 3; channel++)
            samples.at<float>(row*image.cols+col,channel) =
                (uchar) image.at<Vec3b>(row,col)[channel];
// Apply k-means clustering, determining the cluster
// centres and a label for each pixel.
....
```

Regions

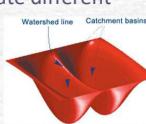
Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 15

15

Watershed segmentation

- In geology watersheds separate different catchment basins.
- Where the rain gets caught
- In computer vision we
 - Identify all minima.
 - Label as different regions.
 - Flood from the minima extending the regions.
 - Where regions meet we have watershed lines.
- Minimum what?
 - Greyscale
 - Gradient
 - Inverse of the chamfer distance



Images from <http://www.mathworks.com/matlabcentral/fileexchange/36380-a-practical-introduction-to-computer-vision-with-opencv>. © 1998-2017 The MathWorks, Inc.

Regions

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 17

17

k means Clustering

```
Mat labels, centres;
kmeans(samples, k, labels, TermCriteria( CV_TERMCRIT_ITER |
    CV_TERMCRIT_EPS, 0.0001, 10000), iterations,
    KMEANS_PP_CENTERS, centres );
// Use centres and label to populate result image
Mat& result_image = Mat( image.size(), image.type() );
for(int row=0; row<image.rows; row++)
    for(int col=0; col<image.cols; col++)
        for (int channel=0; channel < 3; channel++)
            result_image.at<Vec3b>(row,col)[channel] =
                (uchar) centres.at<float>( * (labels.ptr<int>(row*image.cols+col)), channel);
```

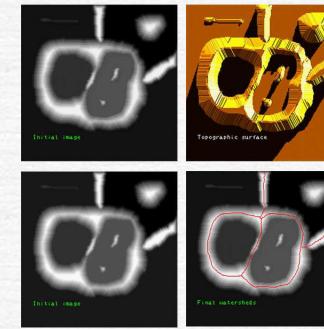
Regions

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 16

16

Watershed segmentation example



Images from <http://www.mathworks.com/help/images/watershed.html>. Reproduced with permission. © 2010, Serge Beucher

Regions

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 18

18

Watershed segmentation with markers

Regions Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014 Slide 19

19

Mean Shift Segmentation

- k-means clustering
 - Requires the number of clusters to be known
 - Takes no account of spatial location
- Mean Shift Segmentation...
 - Do not need to know the number of clusters
 - Can provide spatial as well as colour segmentation
 - Developed by Comaniciu & Meer 2002

Regions Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014 Slide 21

21

Watershed segmentation with markers

- Generally we get too many regions
 - Use a priori labels to identify 'objects' and expand from these rather than the minima.

Regions Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014 Slide 20

20

10

Mean Shift for segmentation– goal

- To associate each point with a particular high-density cluster/mode in colour space.
- Move particles
 - in the direction of
 - the local
 - increasing density

Regions Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014 Slide 22

11

Kernel Density Estimation

- The Problem: Given a sparse dataset determine an estimate of density at each point.

$$\hat{f}_h(x) = \frac{1}{n} \sum_{i=1}^n K_h(x - x_i) = \frac{1}{nh} \sum_{i=1}^n K\left(\frac{x - x_i}{h}\right)$$

- Kernel Density $K()$ function n samples x_i
- h is the bandwidth

- Effectively
 - Smooth all data samples
 - Add them all together
- If the dataset is multidimensional:

$$\hat{f}(\mathbf{x}) = \frac{1}{nh^d} \sum_{i=1}^n K\left(\frac{\mathbf{x} - \mathbf{x}_i}{h}\right).$$

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Regions Slide 23

23

Mean Shift – Basic Algorithm

- For each particle (pixel)
 - Estimate the local kernel density and more importantly the direction of local increasing density (the *mean shift vector*)
 - Shift the particle to the new mean.
 - Re-compute until the location stabilizes.
- Finally identify which pixels ended up in the same location
 - Mark these as members of the same cluster.
 - Determine the local mean of similar particles

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Regions Slide 25

25

Kernel Density Estimation – Kernels

- Many different kernels.
- Kernel functions must integrate to 1.
- Typically use either
 - Uniform

$$k_E(x) = \begin{cases} 1-x & 0 \leq x \leq 1 \\ 0 & x > 1, \end{cases}$$

- Gaussian

$$K_N(\mathbf{x}) = (2\pi)^{-d/2} \exp\left(-\frac{1}{2} \|\mathbf{x}\|^2\right)$$

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Regions Slide 24

24

Mean Shift – The local kernel density

- We limit the points included in the kernel density estimate based on distance and on similar to the current point.

$$K_{h_s, h_r}(\mathbf{x}) = \frac{C}{h_s^d h_r^d} k\left(\frac{\|\mathbf{x}^s\|^2}{h_s}\right) k\left(\frac{\|\mathbf{x}^r\|^2}{h_r}\right)$$

- We must use both a spatial kernel and a colour kernel.
- Both can be Gaussian
- Spatial kernel limits/weights the region to consider around the current point
- Colour/range kernel limits/weights the colour/intensity of the points to be included in the mean.

Maths from "Mean Shift: A Robust Approach Toward Feature Space Analysis" Dorin Comaniciu, and Peter Meer , IEEE TRANSACTIONS ON PATTERN ANALYSIS AND MACHINE INTELLIGENCE, VOL. 24, NO. 5, MAY 2002;

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Regions Slide 26

26

Mean Shift – The mean shift vector

- For radially symmetric kernels $K(\mathbf{x}) = c_{k,d} k(\|\mathbf{x}\|^2)$
- So the KDE becomes $\hat{f}_{h,K}(\mathbf{x}) = \frac{c_{k,d}}{nh^d} \sum_{i=1}^n k\left(\frac{\|\mathbf{x} - \mathbf{x}_i\|^2}{h}\right)$ as $\hat{f}(\mathbf{x}) = \frac{1}{nh^d} \sum_{i=1}^n K\left(\frac{\mathbf{x} - \mathbf{x}_i}{h}\right)$
- We are interested in the rate of change...

$$\hat{\nabla} f_{h,K}(\mathbf{x}) \equiv \nabla \hat{f}_{h,K}(\mathbf{x}) = \frac{2c_{k,d}}{nh^{d+2}} \sum_{i=1}^n (\mathbf{x} - \mathbf{x}_i) k'\left(\frac{\|\mathbf{x} - \mathbf{x}_i\|^2}{h}\right)$$

Set $g(\mathbf{x}) = -k'(\mathbf{x})$ so

$$= \frac{2c_{k,d}}{nh^{d+2}} \sum_{i=1}^n (\mathbf{x}_i - \mathbf{x}) g\left(\frac{\|\mathbf{x} - \mathbf{x}_i\|^2}{h}\right)$$

$$= \frac{2c_{k,d}}{nh^{d+2}} \left[\sum_{i=1}^n g\left(\frac{\|\mathbf{x} - \mathbf{x}_i\|^2}{h}\right) \right] \left[\frac{\sum_{i=1}^n \mathbf{x}_i g\left(\frac{\|\mathbf{x} - \mathbf{x}_i\|^2}{h}\right)}{\sum_{i=1}^n g\left(\frac{\|\mathbf{x} - \mathbf{x}_i\|^2}{h}\right)} - \mathbf{x} \right]$$

Define the mean shift vector as

$$\mathbf{m}_{h,G}(\mathbf{x}) = \frac{\sum_{i=1}^n \mathbf{x}_i g\left(\frac{\|\mathbf{x} - \mathbf{x}_i\|^2}{h}\right)}{\sum_{i=1}^n g\left(\frac{\|\mathbf{x} - \mathbf{x}_i\|^2}{h}\right)} - \mathbf{x}$$

Maths from "Mean Shift: A Robust Approach Toward Feature Space Analysis" Dorin Comaniciu, and Peter Meer, IEEE TRANSACTIONS ON PATTERN ANALYSIS AND MACHINE INTELLIGENCE, VOL. 24, NO. 5, MAY 2002;

Regions

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 27

Mean Shift – Pros and cons

- + Do not need to know the number of clusters *a priori*.
- + Provides spatial as well as colour segmentation.
- Selection of kernel widths can be very hard.
- It is quite slow particularly if there are a lot of clusters.

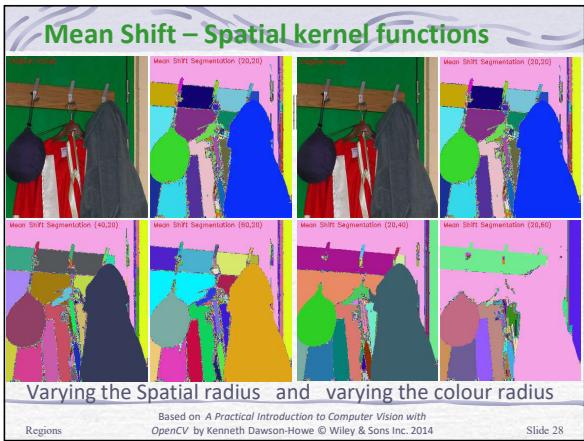


Images from "Mean Shift: A Robust Approach Toward Feature Space Analysis" Dorin Comaniciu, and Peter Meer , IEEE TRANSACTIONS ON PATTERN ANALYSIS AND MACHINE INTELLIGENCE, VOL. 24, NO. 5, MAY 2002; From <http://www.comaniciu.net/Papers/MsRobustApproach.pdf>

Regions

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 29



28

Video

- Introduction
 - Static Background
- Background Models
 - Median
 - Gaussian Mixture Model
 - Shadow Detection
- Tracking
 - Exhaustive Search, Mean Shift, Optical Flow, Feature Point Tracking

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Video Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014 Slide 1

1

Introduction – Common problems

- Illumination & appearance changes
 - Gradual (e.g. time of day)
 - Sudden (e.g. clouds, lights)
 - Shadows
 - Weather (e.g. rain, snow)
- Background changes
 - Objects becoming part of the background
 - Objects leaving the background
 - Background objects oscillating slightly
- Setup
 - Camera motion
 - Frame rate
 - Field of view
 - Distance to objects
 - Location of camera



© Ahmed Elgammal (Reproduced with permission) from <ftp://www.umiacs.umd.edu/pub/elgammal/video/index.html>

Video Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014 Slide 3

3

Introduction – Objects of interest

- Reliably detecting moving objects of interest in a scene.
 1. Motion detection
 2. Moving object detection & location
 3. Derivation of 3D object properties
- When is an object of interest?
 - Size
 - Max and min velocity and acceleration
 - Assumptions:
 - Mutual correspondence
 - Common motion

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Video Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014 Slide 2

2

Introduction – Static background



- Image subtraction $d(i,j) = |f_k(i,j) - b(i,j)|$
 - Difference of current frame and background image..
- `absdiff(frame, background, difference);`
- Background $b(i,j)$ = first frame?

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Video Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014 Slide 4

4

1

2

Introduction – Thresholding

- Threshold result

$$d(i,j) = \begin{cases} 0 & \text{if } |f_k(i,j) - b(i,j)| < T \\ 1 & \text{otherwise} \end{cases}$$

- Threshold T sensitivity?
 - Too high → False negatives
 - Too low → False positives
 - Just right??
- High contrast required.

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Video

Slide 5

5

Introduction – Static Background

```
absdiff( current_frame, first_frame, difference );
cvtColor( difference, moving_points, CV_BGR2GRAY );
threshold( moving_points, moving_points, 30, 255, THRESH_BINARY );
Mat display_image = Mat::zeros( moving_points.size(), CV_8UC3 );
current_frame.copyTo( display_image, moving_points );
```

© Reproduced by permission of Dr. James Ferryman, University of Reading

Video

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 7

7

Introduction – Colour

- But how did we deal with colour?
- Average difference of all channels?
- Just process hue channel?
 - What about white/grey/black?
- Threshold channels separately
 - Any one channel or all channels?
 - Which colour model?

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Video

Slide 6

6

Background Models – Median

- Median: Middle value (from an ordered list)

$$h_n(i,j,p) = \sum_{k=(n-m+1),n} \begin{cases} 1 & \dots \text{if } (f_k(i,j)) = p \\ 0 & \dots \text{otherwise} \end{cases}$$

- No. of frames (m)
- Histogram quantisation?
- Computational expense
 - Adding, storing and removing frames
 - Change in median can be tracked inexpensively from frame to frame
 - Can be approximated using aging

$$h_n(i,j,p) = \sum_{k=1,n} \begin{cases} w_k & \dots \text{if } (f_k(i,j)) = p \\ 0 & \dots \text{otherwise} \end{cases}$$

where $w_1 = 1$ and $w_k = w_{k-1} * 1.001$

- Can also use selective update
- Could use the Mode instead... (Most common value)

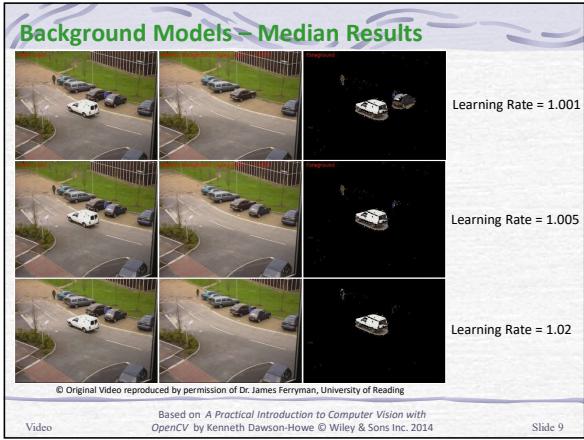
$$b_n(i,j) = p \text{ where } h_n(i,j,p) \geq h_n(i,j,q) \text{ for all } q \neq p$$

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Video

Slide 8

8



9

Background Models – Gaussian Mixture Model

- How to deal with multi-modal background pixels
 - e.g. from trees, water
 - Stauffer & Grimson, 2000
 - Algorithm presented is based on that in Sonka (3rd edition) pp.777-780
- Model multiple values (3-5) at each point.
- Unsupervised learning...
- Most popular method for background modelling

Frame = 1 GMM Background Foreground

© Reproduced by permission of Dr. James Ferryman, University of Reading

Video Based on A Practical Introduction to Computer Vision with OpenCV by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014 Slide 11

11

Background Models – Median Algorithm

First frame:

```

total = 1
for all pixels (i,j)
  median = f1(i,j)
  less_than_median(i,j) = 0
  
```

frame = 1 Median Background foreground

Current frame (n):

```

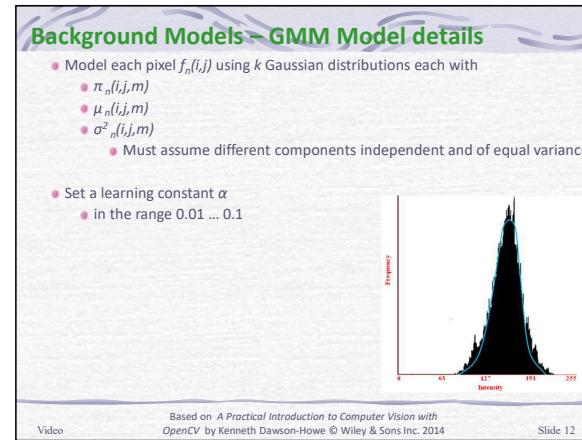
total = total + wn
for all pixels (i,j)
  if (median(i,j) > fn(i,j))
    less_than_median(i,j) = less_than_median(i,j) + wn
  while (less_than_median(i,j) + hn(i,j,median(i,j)) < total/2)
    less_than_median(i,j) = less_than_median(i,j) + hn(i,j,median(i,j))
    median(i,j) = median(i,j) + 1
  while (less_than_median(i,j) > total/2)
    median(i,j) = median(i,j) - 1
    less_than_median(i,j) = less_than_median(i,j) - hn(i,j,median(i,j))
  
```

frame Median Background foreground

Video Based on A Practical Introduction to Computer Vision with OpenCV by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014 Slide 10

10

5



12

6

Background Models – GMM Model Update

- For a new sample $f_n(i,j)$
 - Select the best close Gaussian distribution
 - $\text{close} = \text{within } 2.5 \sigma_n(i,j,m) \text{ of } \mu_n(i,j,m)$
 - If there is a best close Gaussian /

$$\pi_{n+1}(i,j,m) = \alpha * O_n(i,j,m) + (1-\alpha) * \pi_n(i,j,m)$$

where $O_n(i,j,m) = 1$ for the close Gaussian distribution and
0 otherwise

$$\mu_{n+1}(i,j,m) = \mu_n(i,j,m) + O_n(i,j,m) * (\alpha / \pi_{n+1}(i,j,m)) * (f_n(i,j) - \mu_n(i,j,m))$$

$$\sigma^2_{n+1}(i,j,m) = \sigma^2_{n+1}(i,j,m) + O_n(i,j,m) * (\alpha / \pi_{n+1}(i,j,m)) * ((f_n(i,j) - \mu_n(i,j,m))^2 - \sigma^2_n(i,j,m))$$
 - If there is no close Gaussian (replace one...)
 - $x = \arg\min_m \{\pi_n(i,j,m)\}$
 - $\mu_{n+1}(i,j,x) = f_n(i,j)$
 - $\sigma^2_{n+1}(i,j,x) = 2 \cdot \max_m \sigma^2_n(i,j,m)$
 - $\pi_{n+1}(i,j,x) = \frac{1}{2} \cdot \min_m \pi_n(i,j,m)$

Video

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 13

Background Models – GMM issues

GMM issues:

- may fail under fast variations,
- low sensitivity around Gaussian tails,
- less frequent events produce low probability & high variance,
- needs to compute floating point probabilities.

Video

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 15

13

15

Background Models – GMM Moving Points

- Identifying background distributions
 - Define T , a proportion of the frames in which background pixels should be visible.
 - Order the Gaussians by $\pi_{n+1}(i,j,m) / \sigma_{n+1}(i,j,m)$
 - Gaussians 1..B are considered background where

$$B = \arg\min_b \left(\sum_{b=1..m} \pi_{n+1}(i,j,m) \right) > T$$
- Just check if best close Gaussian (or the new Gaussian distribution) is a background distribution
- Finally use morphological dilations and erosions to remove small regions and fill in holes.



© Reproduced by permission of Dr. James Ferryman, University of Reading

Video

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 14

14

Background Models – Shadow Detection

$$SP_k(i,j) = \begin{cases} 1 & \dots \text{if } \left(\alpha < \frac{f_k^V(i,j)}{B_k^V(i,j)} < \beta \right) \text{ and } \left(|f_k^S(i,j) - B_k^S(i,j)| < \tau_S \right) \text{ and } \left(|f_k^H(i,j) - B_k^H(i,j)| < \tau_H \right) \\ 0 & \dots \text{otherwise} \end{cases}$$

- Compare current frame with background image...

(Prati 2003)



- Intensity / luminance / value drops
- Intensity / luminance / value has not dropped too much
- Saturation does not increase too much
- Hue does not change too much

- Hue unpredictable & Change in Luminance can be small (Tattersall 2003)

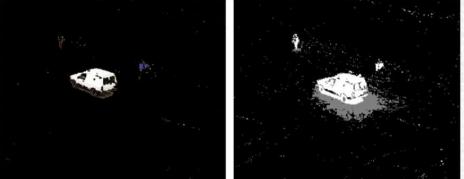
$$SP_k(i,j) = \begin{cases} 1 & \dots \text{if } \left(\lambda < \frac{f_k^V(i,j)}{B_k^V(i,j)} < 1.0 \right) \text{ and } \left(|f_k^S(i,j) - B_k^S(i,j)| < \tau_S \right) \\ 0 & \dots \text{otherwise} \end{cases}$$

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 16

16

Background Models – Shadow Detection



© Reproduced by permission of Dr. James Ferryman, University of Reading

Video Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014 Slide 17

17

Tracking – Introduction

- Used in video surveillance, sports video analysis, vehicle guidance systems, etc.

- A hard task because objects
 - may be undergoing complex motion
 - may change shape
 - may be occluded
 - may change appearance due to lighting/weather
 - may physically change appearance

Video Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014 Slide 19

19

Background Models – Shadow Detection – Code

```
Ptr<BackgroundSubtractorMOG2> gmm =
    createBackgroundSubtractorMOG2();
gmm->apply(current_frame, foreground_mask);
gmm( current_frame, foreground_mask );
threshold( foreground_mask, moving_points, 150, 255, THRESH_BINARY );
...
threshold( foreground_mask, changing_points, 50, 255, THRESH_BINARY );
absdiff( moving_points, changing_points, shadow_points );
...
Mat mean_background_image;
gmm->getBackgroundImage( mean_background_image );
```

Video Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014 Slide 18

18

Tracking – Topics

- Exhaustive search
- Mean Shift
- Optical Flow
- Feature based tracking

Video Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014 Slide 20

20

Tracking – Exhaustive Search

- Extract object to be tracked from frame
- Compare in all possible positions in future frame(s)
 - Use a similarity metric
 - E.g. normalised cross correlation
 - Just pick the best match?
- Need extra degrees of freedom for scale and orientation.
- May fail if object motion is too complex.
- Template matching and chamfer matching support this type of tracking.

Video

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

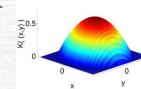
Slide 21

21

Tracking – Mean Shift

- Model the “object”

$$\hat{q}_u \triangleq C \sum_{i=1}^n k \left(\frac{\|\mathbf{x}_i\|^2}{h_q} \right) \delta [b(\mathbf{x}_i) - u]$$
 - Probability Density Function (histogram) of colours.
 - Limit the number of bins
 - Typically an elliptical region is used.
- Weight the values relative to their location.
 - Epanechnikov kernel



Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 23

23

Tracking – Mean Shift

- Tracks “objects” in video by
 - Searching *locally* for the most similar region.
 - Using a histogram to represent the “object”.
 - Using (iterative) gradient ascent to locate the best match.



Video

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

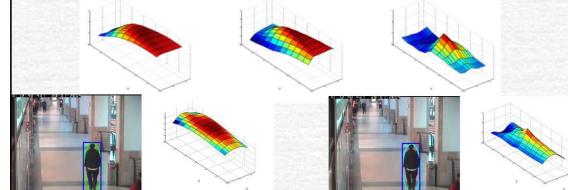
Slide 22

22

Tracking – Mean Shift

- Model candidates regions...

$$\hat{p}_u(\mathbf{y}) \triangleq C_h \sum_{i=1}^{n_h} k \left(\frac{\|\mathbf{y} - \mathbf{x}_i\|^2}{h_p} \right) \delta [b(\mathbf{x}_i) - u]$$
- Matching to find the best new location...
 - Compare distributions directly.
 - Move to the mode in matching space.
 - Bhattacharya EMD NCC



Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 24

24

Tracking – Mean Shift

Mean shift approach

- Consider the gradient of the similarity function (the Bhattacharyya coefficient)...
- Gradient of superposition of kernels, centered at each data point is equivalent to convolving the data points with the gradient of the kernel

$$w_i = \sum_{u=1}^m \sqrt{\frac{q_u}{p_u(\mathbf{y}_0)}} \delta [b(x_i) - u] \quad \mathbf{y}_1 = \frac{\sum_{i=1}^{n_h} \mathbf{x}_i w_i}{\sum_{i=1}^{n_h} w_i}$$

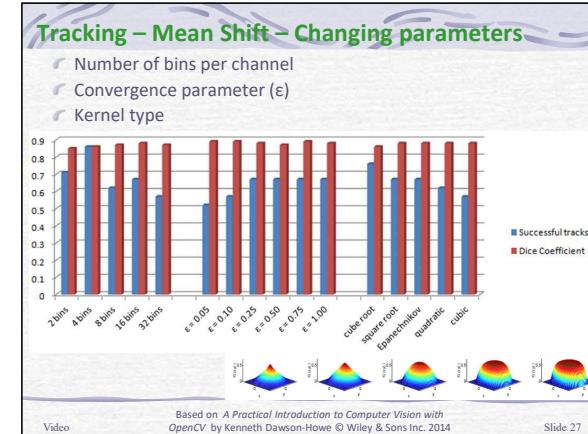
- Derived from the Bhattacharyya similarity measure
- Assumes Epanechnikov kernels
- Move in the direction of the highest gradient

- Iterate until convergence
 - Separation between \mathbf{y}_0 and \mathbf{y}_1 less than ϵ

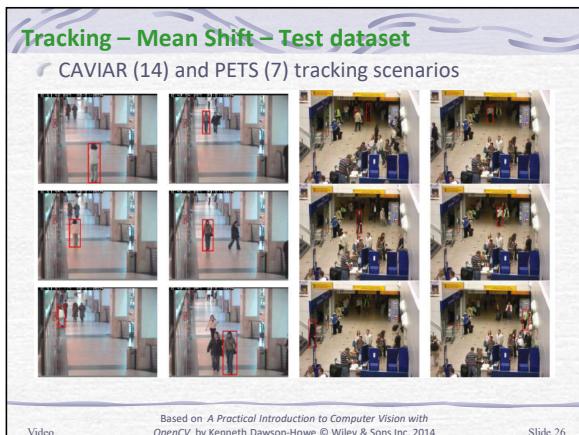
Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Video Slide 25

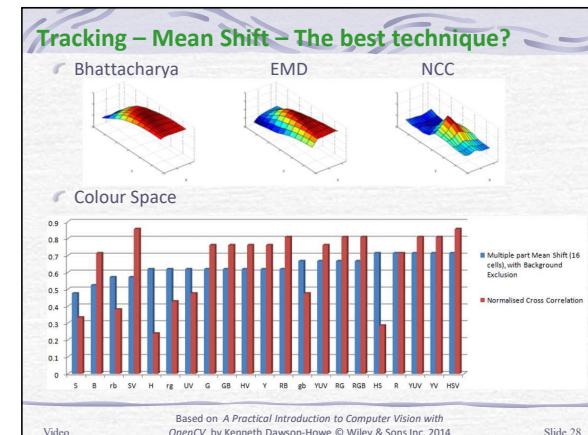
25



27



26



28

Tracking – Mean Shift – Background Exclusion

- IF a background model of the scene is available
 - Favour image regions which are similar to the object model
 - AND dissimilar to the corresponding background region.

Analogous to background subtraction
No image differencing.

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Video Slide 29

29

Tracking – Mean Shift – In OpenCV

- Back-project a histogram of the object into the current frame
- Searches for a region of the same size within the back projection looking for the highest (weighted) sum.

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Video Slide 31

31

Tracking – Mean Shift – Multipart models

- Histograms have a lack of spatial structure
- To deal with this we can use a multipart model:

Improvements:

Model	Successful Tracks (with background exclusion)	Successful Tracks (without background exclusion)	Median of medians dice coefficient (with background exclusion)	Median of medians dice coefficient (without background exclusion)
Basic mean shift	~0.65	~0.45	~0.85	~0.80
Multipart (left-and-right)	~0.60	~0.40	~0.85	~0.80
Multipart (top-and-bottom)	~0.65	~0.45	~0.85	~0.80
Multipart (bottom)	~0.55	~0.35	~0.85	~0.80
Multipart (top)	~0.60	~0.40	~0.85	~0.80
Multipart (3x3)	~0.70	~0.50	~0.85	~0.80
Multipart (2x2)	~0.75	~0.60	~0.85	~0.80

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Video Slide 30

30

15

Tracking – Mean Shift – In OpenCV

```
Rect position(starting_position);
TermCriteria criteria( cv::TermCriteria::MAX_ITER, 5, 0.01 );
meanShift( back_projection_probabilities, position, criteria);
```

Starting position Chosen channel (blue)
Back projection Current position

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Video Slide 32

32

16

Tracking – Dense Optical Flow

- Compute a motion field (known as optical flow) for the entire image
 - Direction & Magnitude

Frame = 4! Farneback Optical Flow
Final Year Project by Gavin Corkery

© Reproduced by permission of Dr. James Ferryman, University of Reading

Video Based on A Practical Introduction to Computer Vision with OpenCV by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014 Slide 33

33

Tracking – Dense Optical Flow

- To compute the optical flow $(\frac{\Delta i}{\Delta t}, \frac{\Delta j}{\Delta t})$, assuming that the displacement is small...

$$f_{t+\Delta t}(i + \Delta i, j + \Delta j) = f_t(i, j) + \frac{\partial f}{\partial t} \Delta t + \frac{\partial f}{\partial i} \Delta i + \frac{\partial f}{\partial j} \Delta j$$
- Hence (given our previous equation)

$$\frac{\partial f}{\partial t} \Delta t + \frac{\partial f}{\partial i} \Delta i + \frac{\partial f}{\partial j} \Delta j = 0$$
- So

$$\frac{\partial f}{\partial t} \frac{\Delta t}{\Delta t} + \frac{\partial f}{\partial i} \frac{\Delta i}{\Delta t} + \frac{\partial f}{\partial j} \frac{\Delta j}{\Delta t} = 0$$
- And reorganising

$$\left[\frac{\partial f}{\partial i} \frac{\partial f}{\partial j} \right] \begin{bmatrix} \frac{\Delta i}{\Delta t} \\ \frac{\Delta j}{\Delta t} \end{bmatrix} = -\frac{\partial f}{\partial t}$$

Video Based on A Practical Introduction to Computer Vision with OpenCV by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014 Slide 35

35

Tracking – Dense Optical Flow

- Based on the brightness constancy constraint
 - Object points will have the same brightness over a short period of time $f_t(i, j) = f_{t+\Delta t}(i + \Delta i, j + \Delta j)$
- Need to find the displacement $(\Delta i, \Delta j)$ which will minimise the residual error

$$\epsilon(\Delta i, \Delta j) = \sum_{i=i_{current}-w}^{i_{current}+w} \sum_{j=j_{current}-w}^{j_{current}+w} f_t(i, j) - f_{t+\Delta t}(i + \Delta i, j + \Delta j)$$

Video Based on A Practical Introduction to Computer Vision with OpenCV by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014 Slide 34

34

Tracking – Dense Optical Flow – Code

```
calcOpticalFlowFarneback (previous_gray_frame, gray_frame,
                           optical_flow, 0.5, 3, 15, 3, 5, 1.2, 0);
cvtColor (previous_gray_frame, display, CV_GRAY2BGR);
for (int row = 4; row < display.rows; row+=8)
  for(int column = 4; column < display.cols; column+8)
  {
    Point2f& flow = optical_flow.at<Point2f>(row,column);
    line (display, Point(column,row), Point(
      cvRound(column+flow.x), cvRound(row+flow.y)),
      passed_colour);
  }
gmm.getBackgroundImage( mean_background_image );
```

Video Based on A Practical Introduction to Computer Vision with OpenCV by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014 Slide 36

36

Tracking – Dense Optical Flow – Problems

- Optical flow is the apparent motion of points within a scene.
 - Based on brightness patterns
 - Needs texture

Q. What happens if the brightness changes?

- i.e. brightness constancy does not hold.

A. Perhaps look at optical flow in gradient space.

Q. What if a point moves differently to all its neighbours?

A. Use Region based optical flow

Video

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 37

Tracking – Feature-based Optical Flow

- We cannot accurately compute optical flow for constant regions or along edges.
- Often better to compute optical flow just for features... (e.g. Lucas Kanade feature tracker)...



© Reproduced by permission of Dr. James Ferryman, University of Reading

Video

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 39

37

39

Tracking – Dense Optical Flow – Problems

Q. What happens to a rotating sphere? What about a “barber pole”?

A. We get the wrong motion. i.e. it fails.

Q. What happens if the motion is too large?

A. Use Iterative Refinement (Lucas-Kanade)

Video

Based on *A Practical Introduction to Computer Vision with OpenCV* by Kenneth Dawson-Howe © Wiley & Sons Inc. 2014

Slide 38

38

19

20